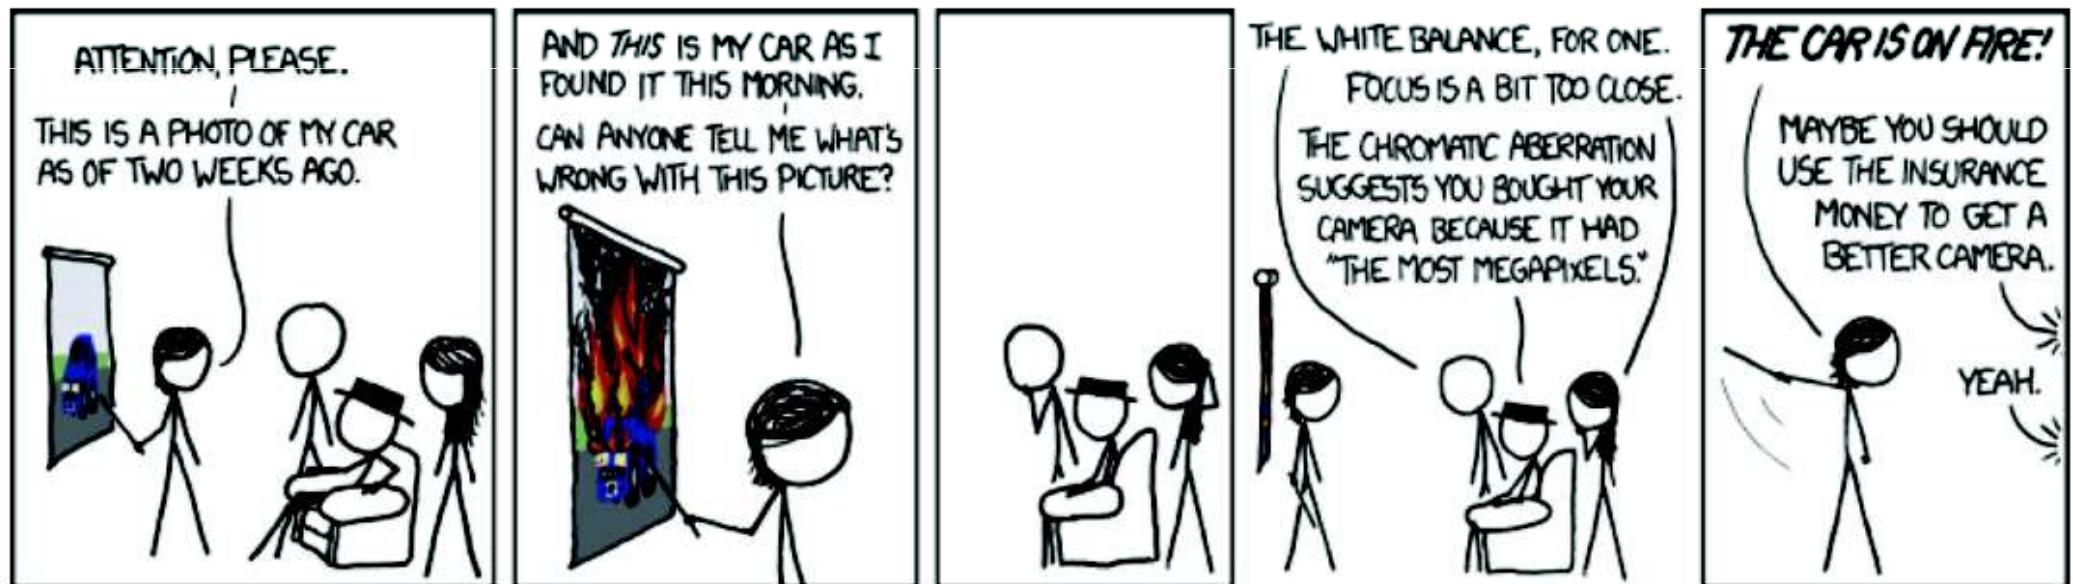


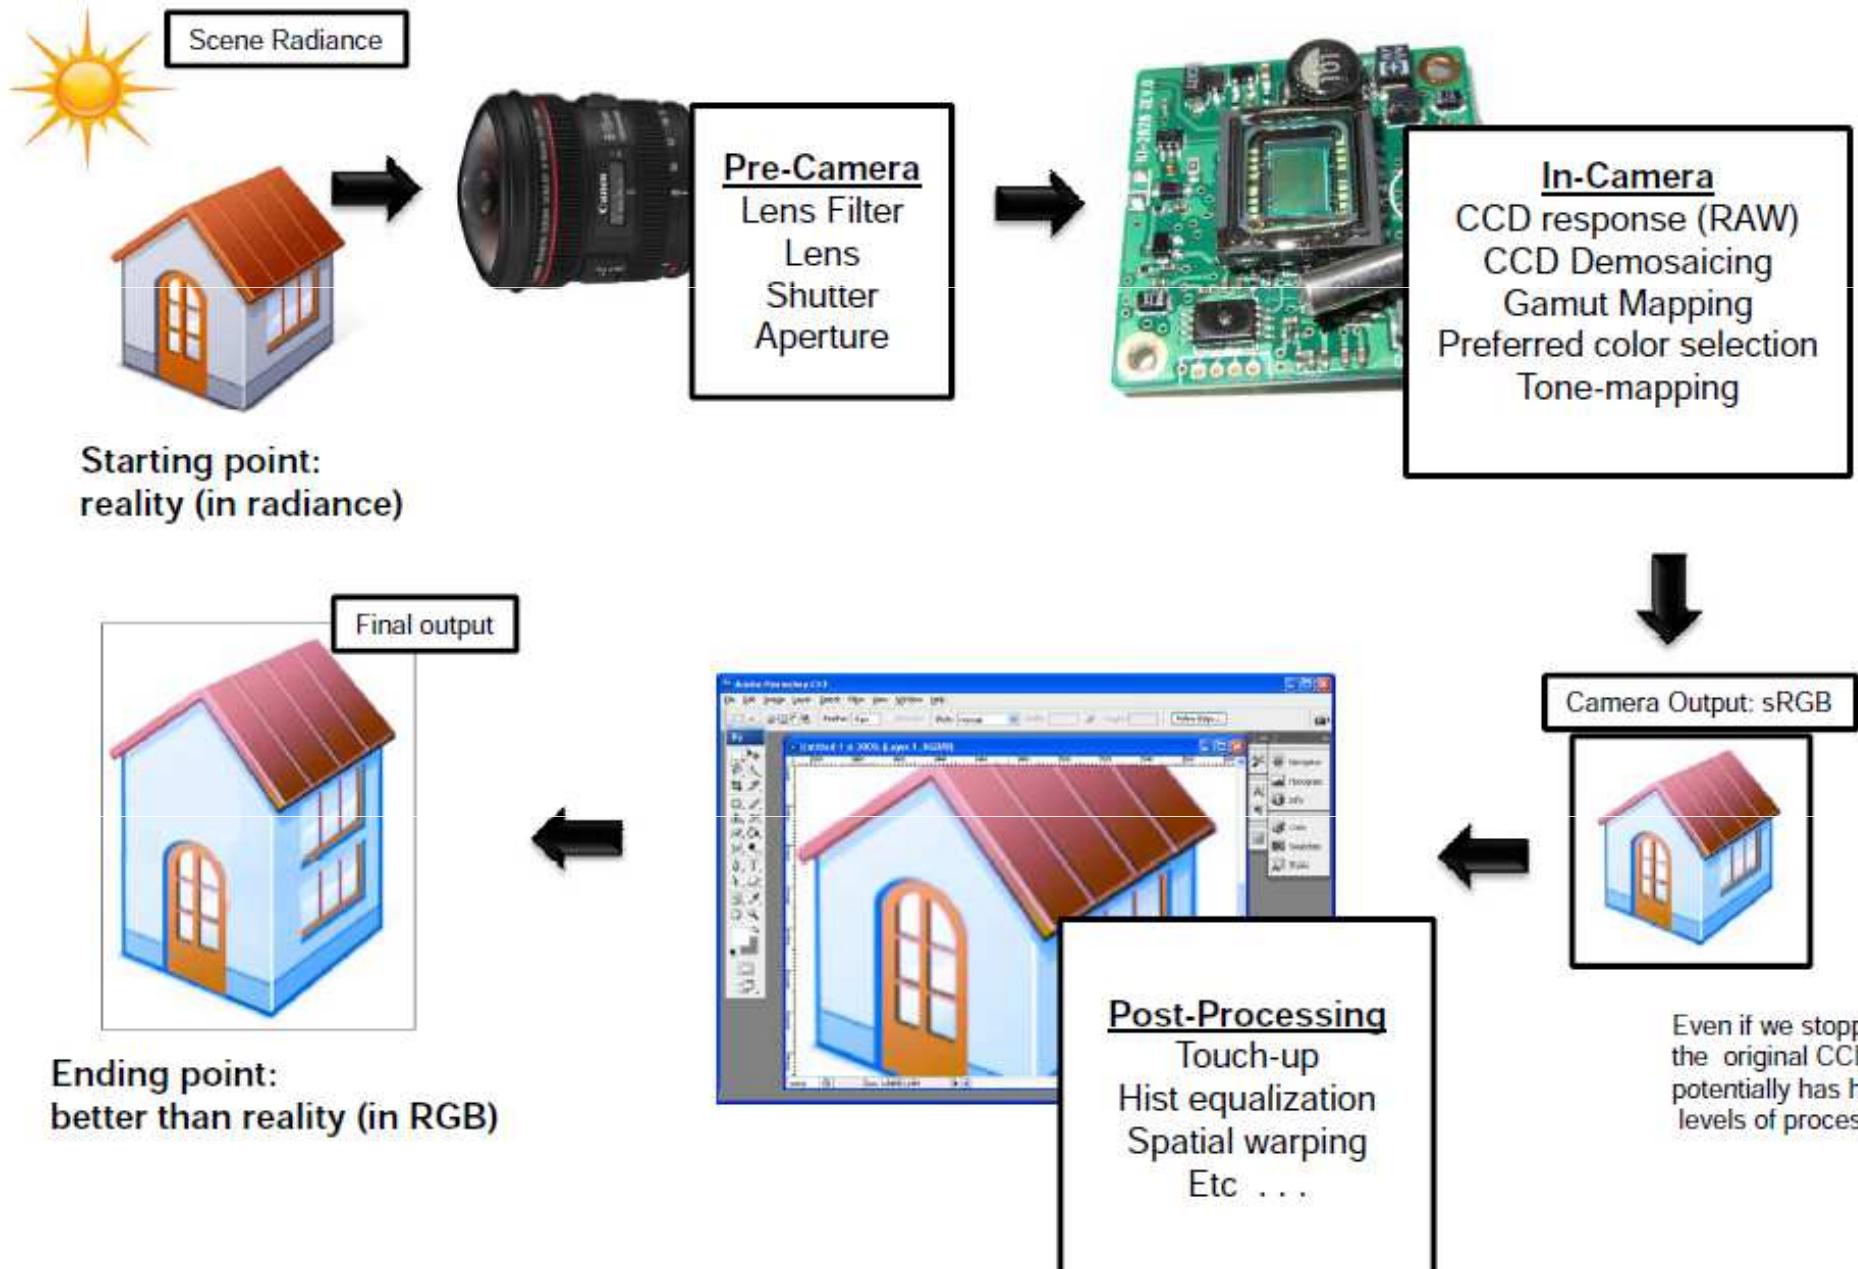
>>> IMAGE PROCESSING AND
COMPUTATIONAL PHOTOGRAPHY
SESSION 2: Basics of image processing

Oriol Pujol & Simone Balocco



<http://xkcd.com/1014/>

Modern image acquisition pipeline

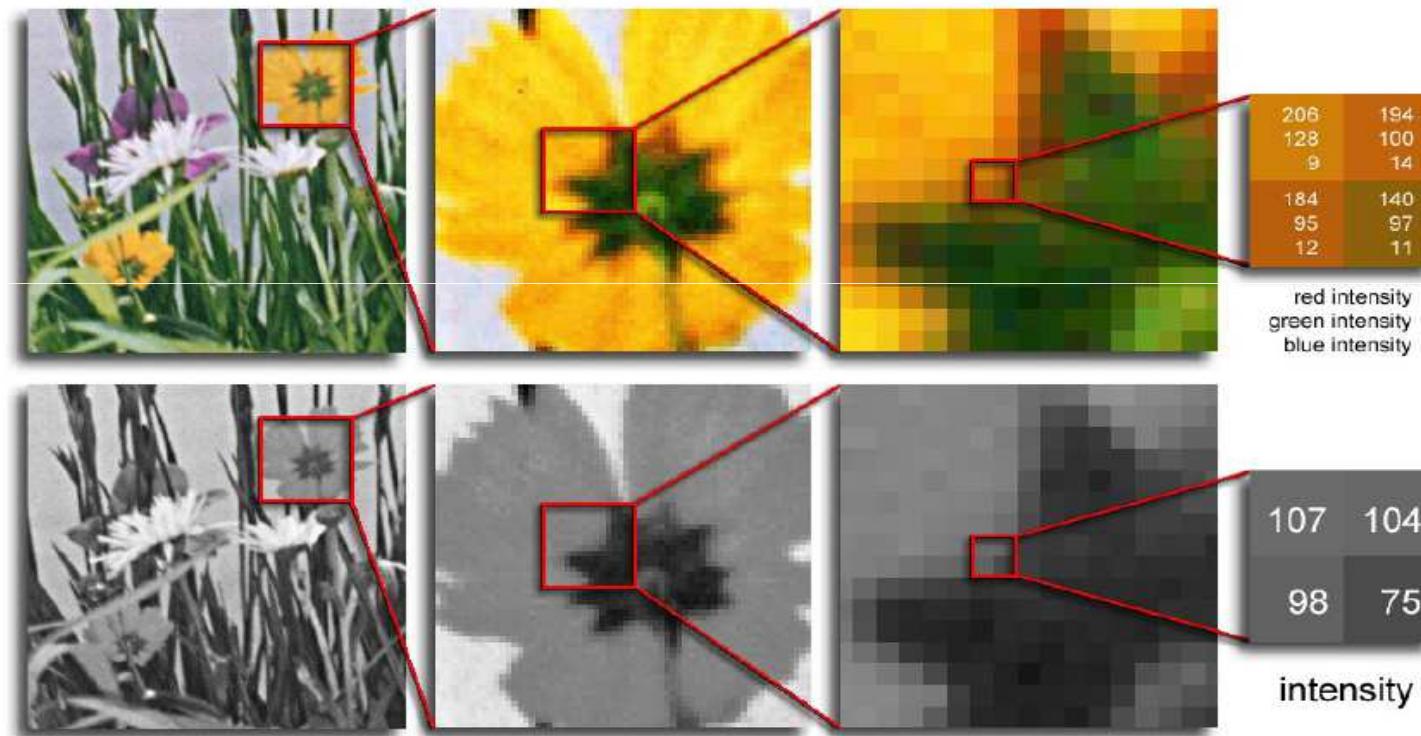


Basics of image processing

Image representation

An image is defined as a 2-dimensional function $f(x,y)$ where x and y are *spatial* coordinates, and the amplitude of f is called *intensity* or *gray level*.

When x, y and f are finite and discrete quantities, we call the image a *digital image*.



Digitalization

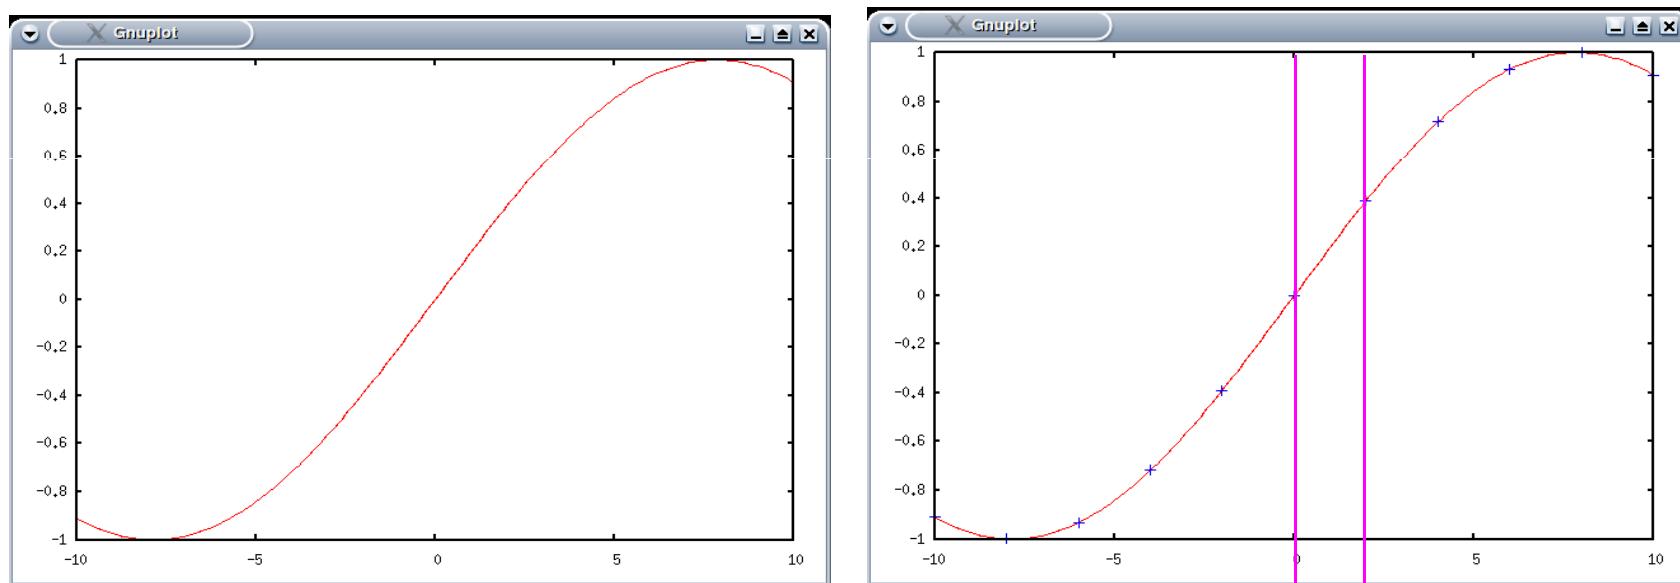
Real data is defined in a continuous domain. However, computers can only handle discrete sets of values.

The way to represent continuous data is to **dicreteize and quantify**.

Discretization

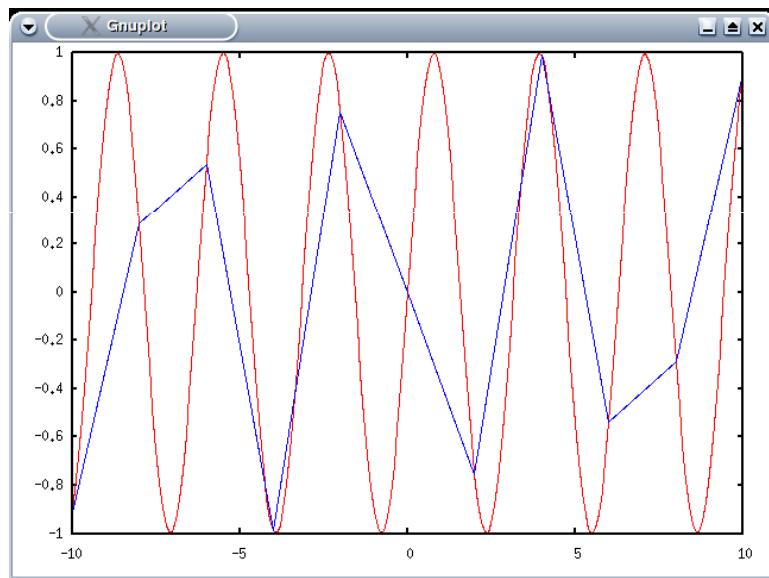
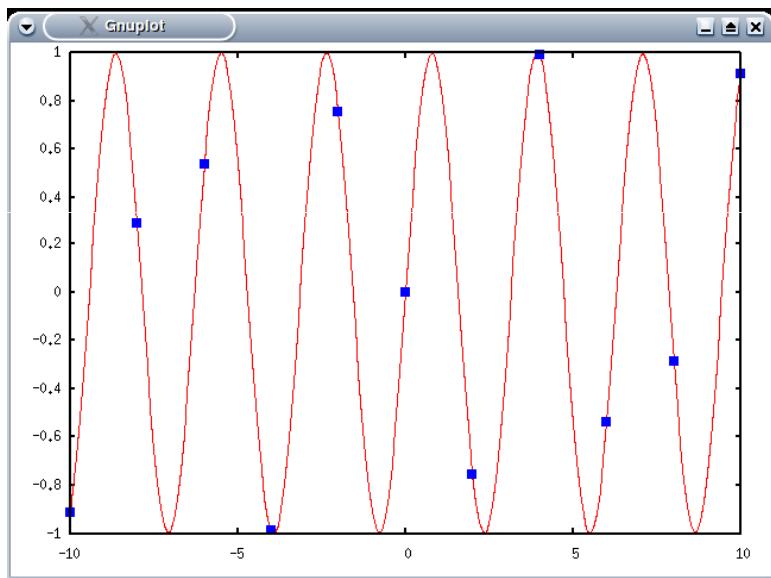
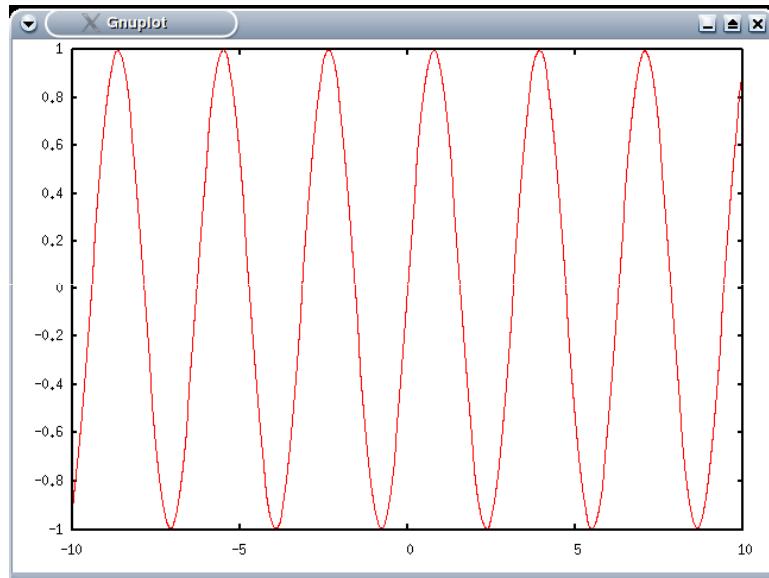
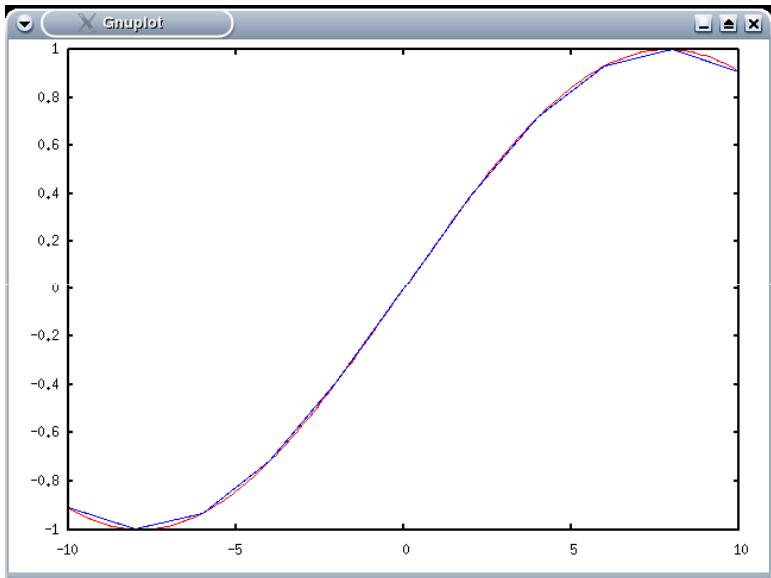
Associated to the concept of sampling and characterized by the sampling frequency or sampling period.

$$x(t) \rightarrow x(n \cdot T_s), \quad \forall n \in \{1 \dots T\}$$



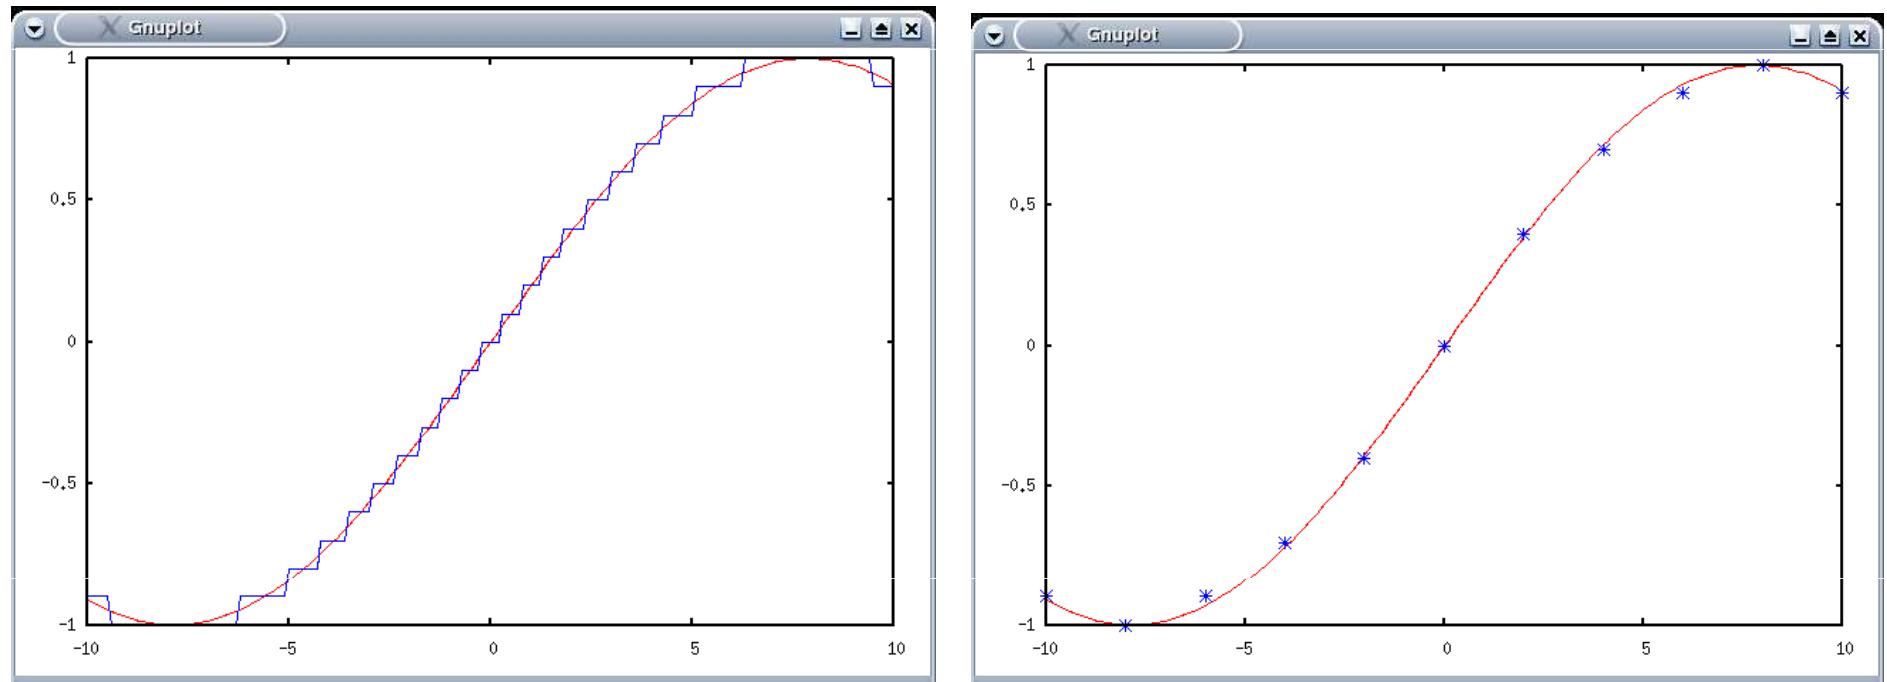
T_s

Reconstruction from sampling



Quantization

Up to this point we have a discrete signal in space/time but the amplitude of the signal is still continuous. The process of discretize the amplitude of the signal is known as **quantization**.



Observe that each level of the quantization correspond to the value stored. In this sense, in an image domain the quantized value is the pixel value.

2D sampling and quantization

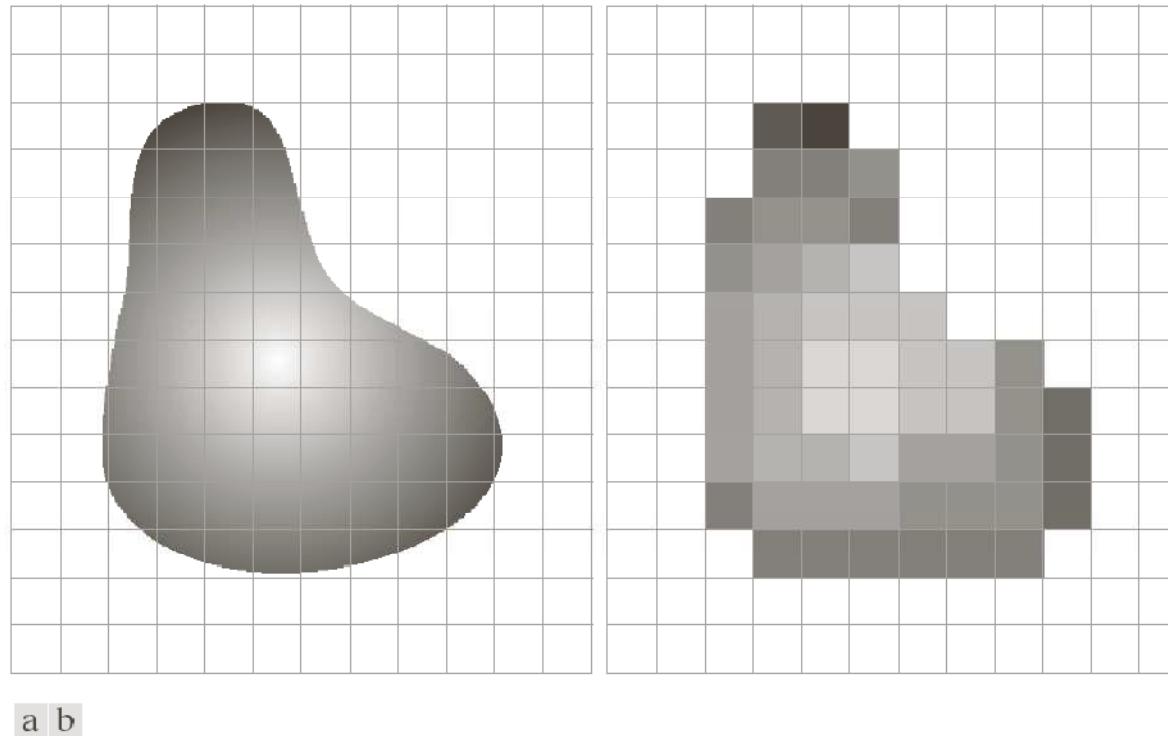


FIGURE 2.17 (a) Continuous image projected onto a sensor array. (b) Result of image sampling and quantization.

The spatial sampling is related to image resolution (e.g. 640x480).
And quantization is related to intensity or color depth (the number of bits per pixel).

Resolution

512x512



64x64



128x128



32x32



Quantization

The number of gray levels associated to b bits is

$$L = 2^b$$



64 levels



16 levels



4 levels



2 levels

Quantization and indexing

The quantization value does not necessarily need to be the extremes of the quantization interval. One may use the mean value or another indexed value. This method is called indexing.



64 levels



16 levels



4 levels



2 levels

Dithering

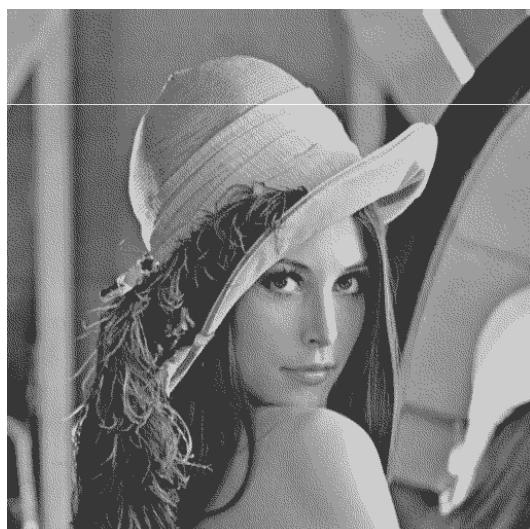
Dither is an intentionally applied form of noise used to randomize quantization error, preventing large-scale patterns such as "banding" ...



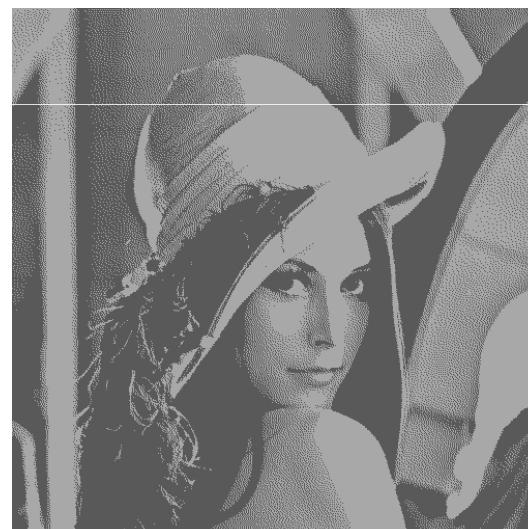
64 levels



16 levels



4 levels



2 levels

Comparison of techniques applied to 2 bits depth images



Brightness

It is a measure of the luminance (global light intensity value) of the image

An estimation of the brightness value:

$$B = \frac{1}{N} \sum_{x,y} f(x, y)$$

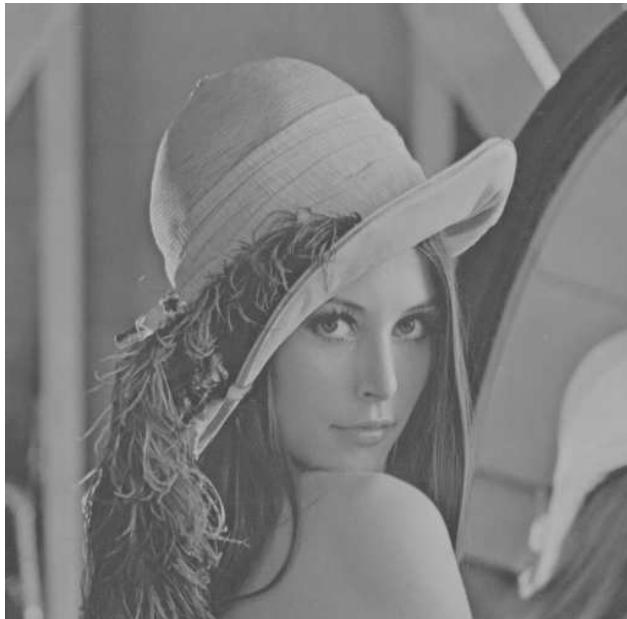


Contrast

Measure of the perceived intensity difference between two regions.

A simple definition of contrast:

$$C = \frac{|f_a - f_b|}{f_a + f_b}$$



Sharpness and Acutance

Acutance is the edge contrast of the image



Sharpness describes the clarity of detail in an image and depends on resolution and acutance



Low res & High ac



High res & Low ac



High res & High ac

Noise

Different types of noise usually categorized by their model:
Additive, multiplicative, impulsive, ...

Examples:

Additive White Gaussian Noise

$$f(x, y) = I(x, y) + n, \quad n \in \mathcal{N}(0, \sigma)$$

Salt and Pepper Noise (Random impulses of extreme values)



Gaussian noise with different sigma

Salt and Pepper

Spatial Domain operations

The term spatial domain refers to the plane of the image.

Spatial methods manipulate pixels themselves.

Basic operation on stand-alone pixels

$$g(x, y) = T[f(x, y)]$$

Operations on pixels based on general image information

$$g(x, y) = T[f(x, y), \theta]$$

Operations taking into account the neighborhood

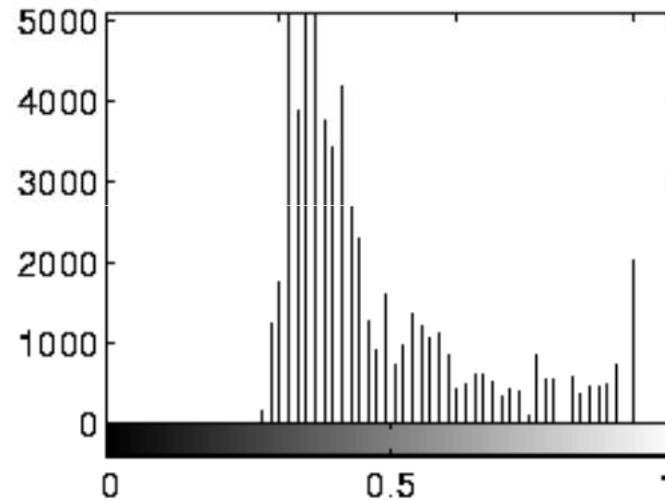
$$g(x, y) = T[f(x, y), \mathcal{N}(x, y)]$$

Histogram equalization

original



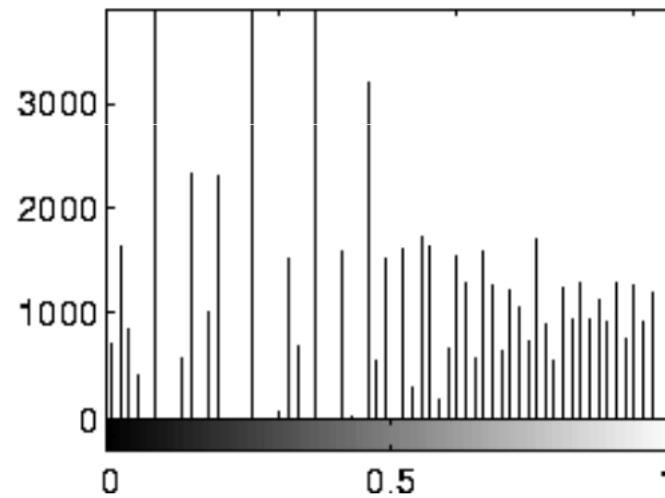
histogram



equalized image



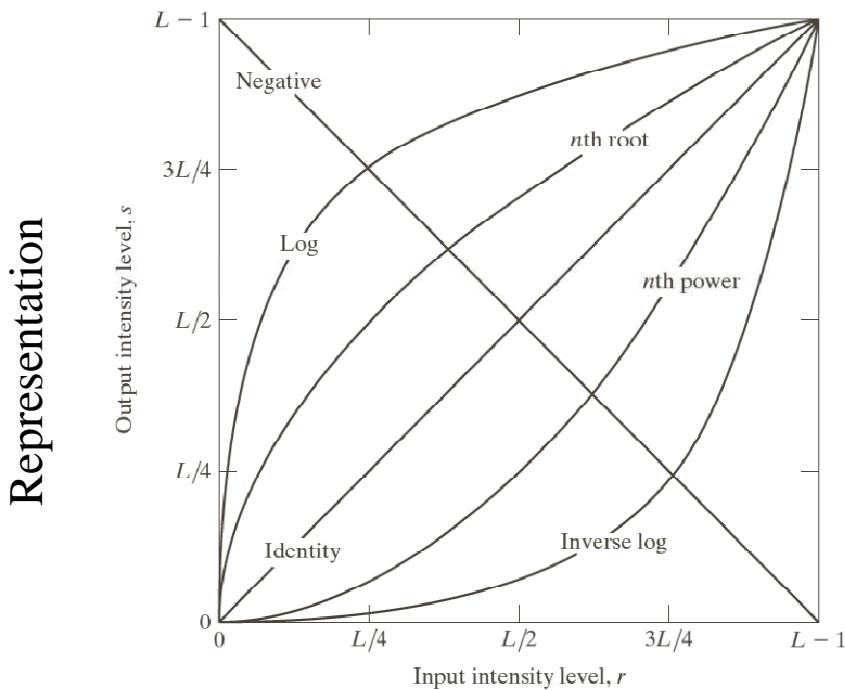
equalized histogram



Basic Intensity Transformation Functions

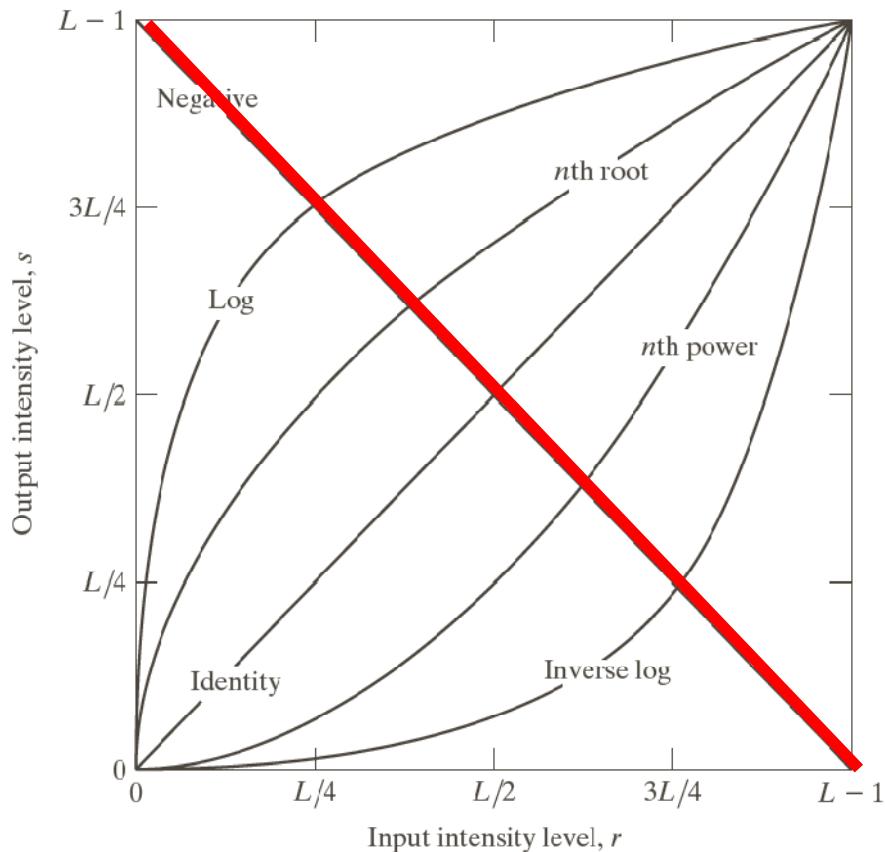
Idea: Replace one intensity value with another.

The simplest transformation is implemented using a Look Up Table (LUT).
For an 8-bit gray image we have to store 256 values.



Negative

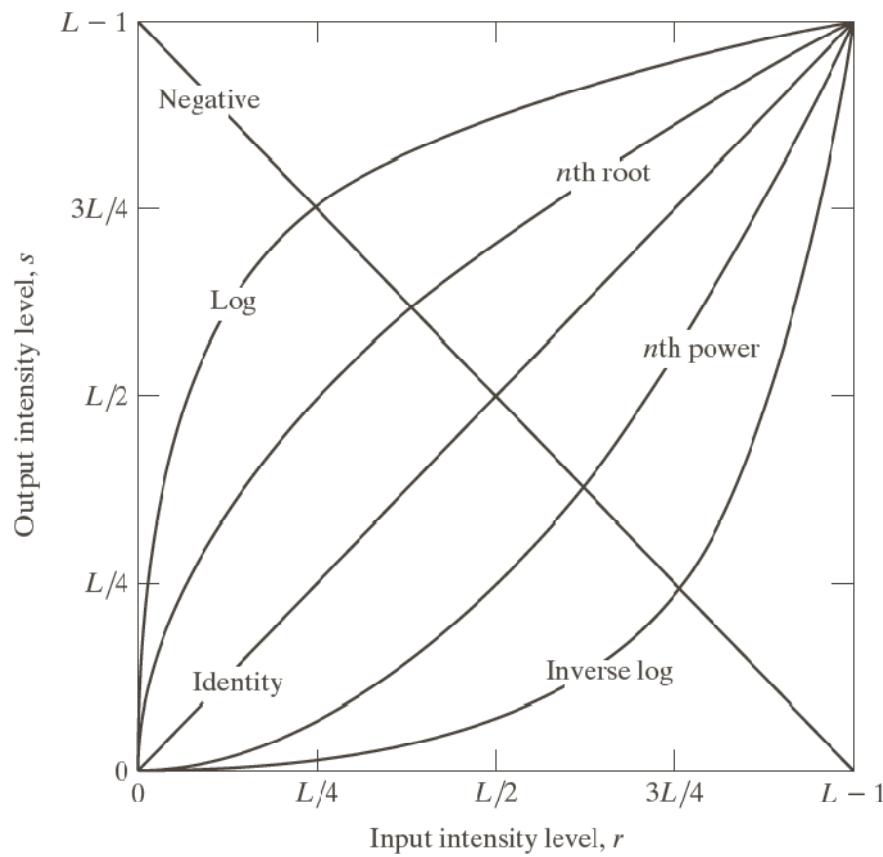
$$g(x, y) = \max(f(x, y)) - f(x, y)$$



Log and Power Transformations

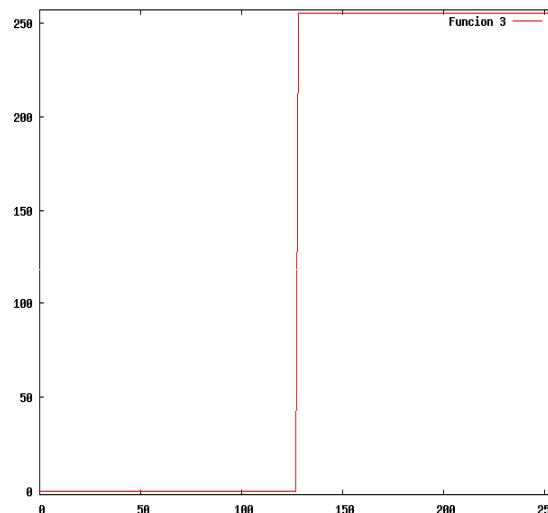
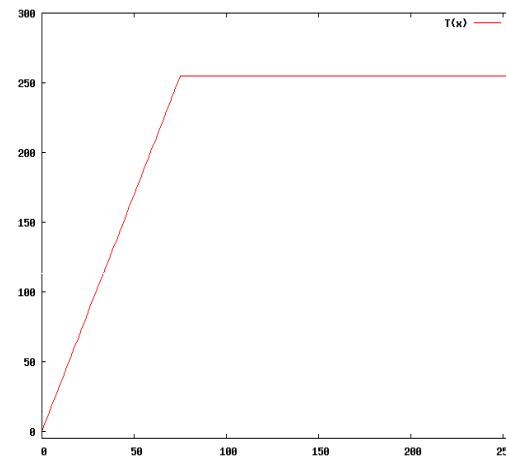
$$g(x, y) = c \log(1 + f(x, y))$$

$$g(x, y) = cf^\gamma(x, y)$$



Effect: compress or expand intensity value range

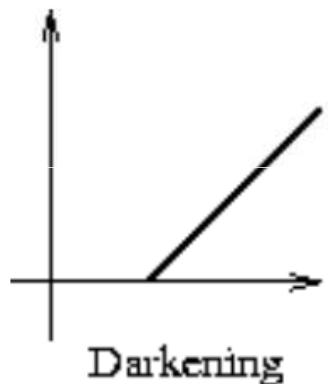
Piece-wise Transformations



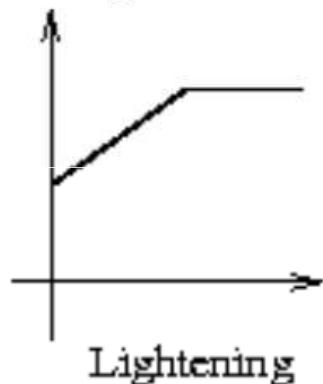
Exercise

Write the transformations or LUTs

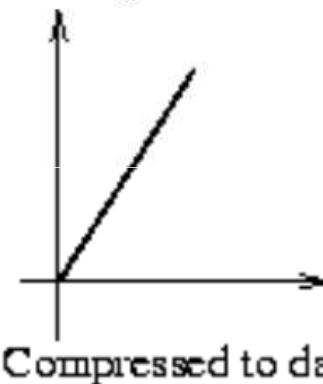
Spatial Domain operations



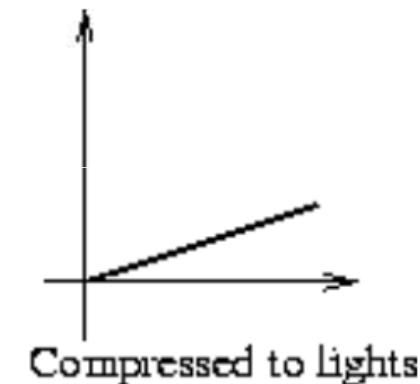
Darkening



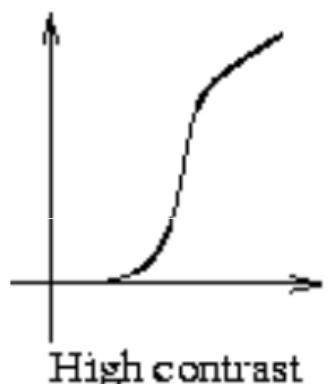
Lightening



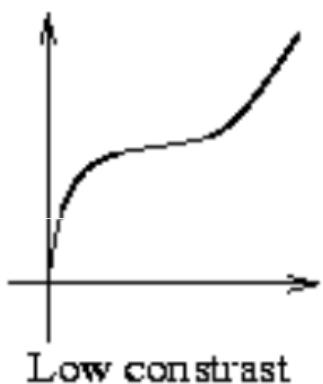
Compressed to darks



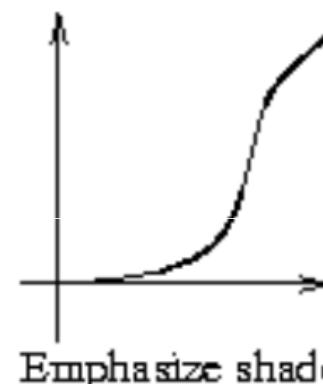
Compressed to lights



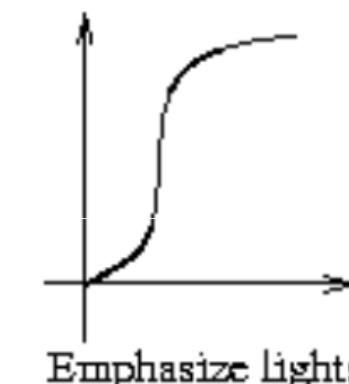
High contrast



Low contrast



Emphasize shadows



Emphasize lights

Introduction

The term spatial domain refers to the plane of the image.

Spatial methods manipulate pixels themselves.

Basic operation on stand-alone pixels

$$g(x, y) = T[f(x, y)]$$

Operations on pixels based on general image information

$$g(x, y) = T[f(x, y), \theta]$$

Operations taking into account the neighborhood

$$g(x, y) = T[f(x, y), \mathcal{N}(x, y)]$$

INTUITIVE MECHANIC OF SPATIAL FILTERING

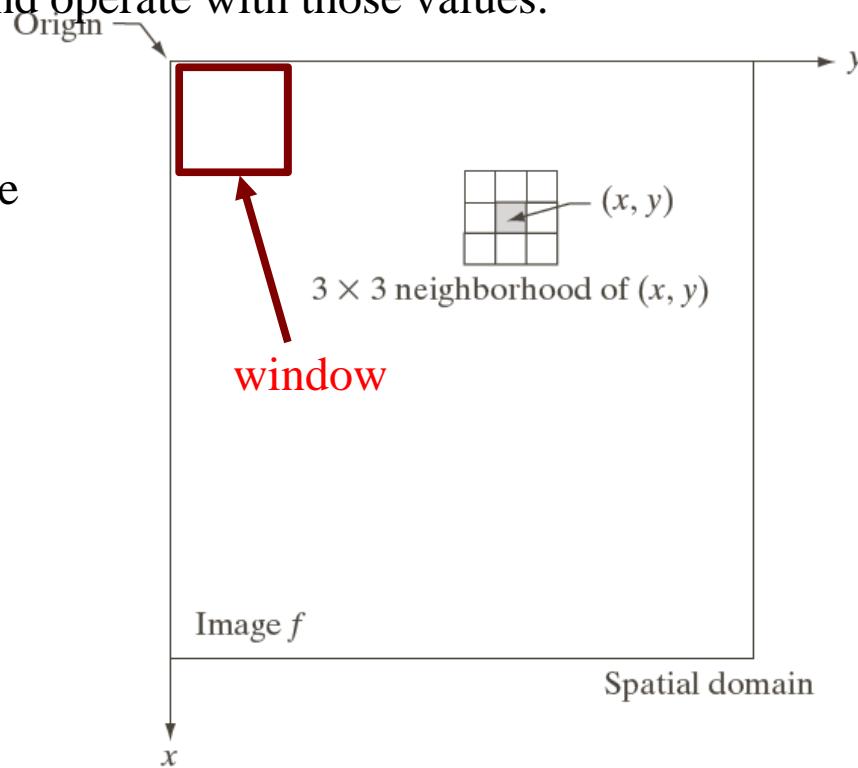
We borrow the name *filtering* from the concepts of frequency filtering, where filtering refers to accepting (passing) or rejecting certain frequency components. [More to come in a few weeks]

So, what must we do if we want to process a pixel according to its neighborhood?

Lay a window around the pixel of interest and operate with those values.

And what do we do if we want to process the full image?

Repeat this process for each pixel



INTUITIVE MECHANIC OF SPATIAL FILTERING

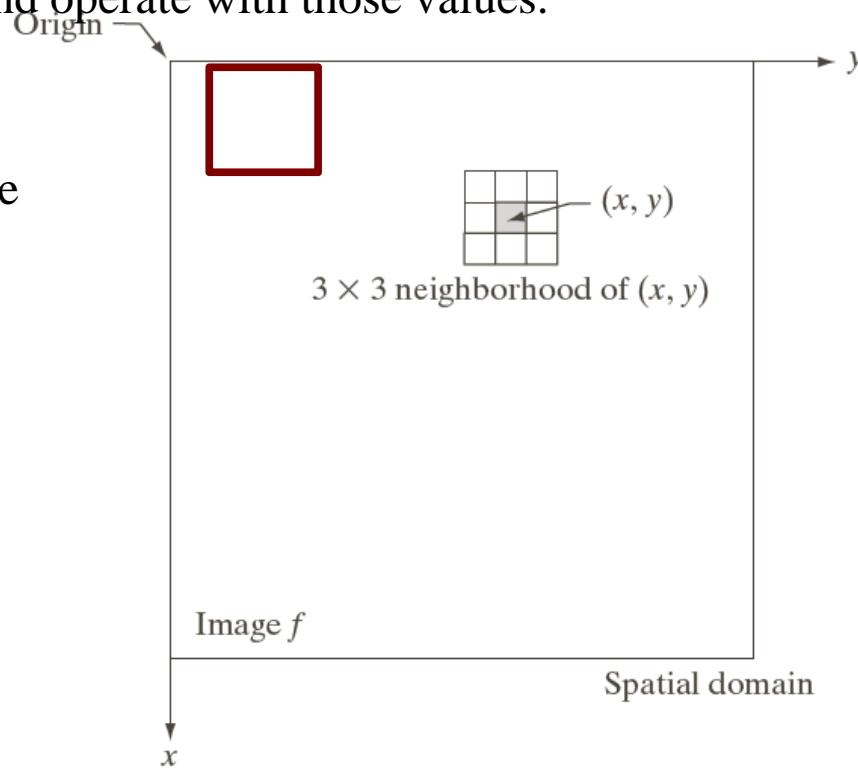
We borrow the name *filtering* from the concepts of frequency filtering, where filtering refers to accepting (passing) or rejecting certain frequency components. [More to come in a few weeks]

So, what must we do if we want to process a pixel according to its neighborhood?

Lay a window around the pixel of interest and operate with those values.

And what do we do if we want to process the full image?

Repeat this process for each pixel



INTUITIVE MECHANIC OF SPATIAL FILTERING

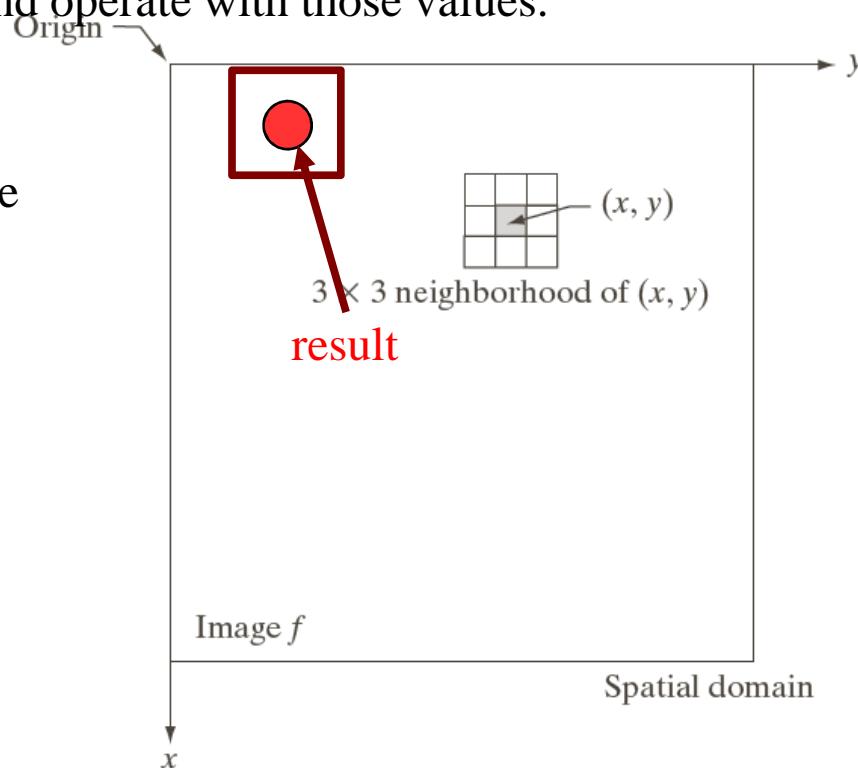
We borrow the name *filtering* from the concepts of frequency filtering, where filtering refers to accepting (passing) or rejecting certain frequency components. [More to come in a few weeks]

So, what must we do if we want to process a pixel according to its neighborhood?

Lay a window around the pixel of interest and operate with those values.

And what do we do if we want to process the full image?

Repeat this process for each pixel



What for?

Enhance images

Denoise

Resize

Increase contrast

Extract features

Texture

Edges

Salient points

Detect patterns

Template

matching

A SIMPLE EXAMPLE

We want to smooth an image. For this task we may use the mean value in a 3x3 window.

$f(x-1,y-1)$	$f(x-1,y)$	$f(x-1,y+1)$
$f(x,y-1)$	$f(x,y)$	$f(x,y+1)$
$f(x+1,y-1)$	$f(x+1,y)$	$f(x+1,y+1)$

$$\begin{aligned} & \frac{1}{9}(f(x-1,y-1) + f(x-1,y) + f(x-1,y+1) + \dots \\ & \quad + f(x,y-1) + f(x,y) + f(x,y+1) + \dots \\ & \quad + f(x+1,y-1) + f(x+1,y) + f(x+1,y+1)) \end{aligned}$$

A SIMPLE EXAMPLE

We want to smooth an image. For this task we may use the mean value in a 3x3 window.

$f(x-1,y-1)$	$f(x-1,y)$	$f(x-1,y+1)$
$f(x,y-1)$	$f(x,y)$	$f(x,y+1)$
$f(x+1,y-1)$	$f(x+1,y)$	$f(x+1,y+1)$

**OBSERVE THAT THE SAME RESULT CAN BE OBTAINED USING
A WINDOW WITH WEIGHTS AND ADDING THE ELEMENT
MULTIPLICATION**

$w(-1,-1)$	$w(-1,0)$	$w(-1,1)$
$w(0,-1)$	$w(0,0)$	$w(0,1)$
$w(1,-1)$	$w(1,0)$	$w(1,1)$

MASK

$$\frac{1}{9} \sum_{i,j \in \{-1,0,1\}} w(i,j) f(x+i, y+j)$$

A SIMPLE EXAMPLE

We want to smooth an image. For this task we may use the mean value in a 3x3 window.

$f(x-1,y-1)$	$f(x-1,y)$	$f(x-1,y+1)$
$f(x,y-1)$	$f(x,y)$	$f(x,y+1)$
$f(x+1,y-1)$	$f(x+1,y)$	$f(x+1,y+1)$

OBSERVE THAT THE SAME RESULT CAN BE OBTAINED USING
A WINDOW WITH WEIGHTS AND ADDING THE ELEMENT
MULTIPLICATION

$w(-1,-1)$	$w(-1,0)$	$w(-1,1)$
$w(0,-1)$	$w(0,0)$	$w(0,1)$
$w(1,-1)$	$w(1,0)$	$w(1,1)$

} MASK / Kernel

WHY 9? $\frac{1}{9} \sum_{i,j \in \{-1,0,1\}} w(i,j) f(x+i, y+j)$

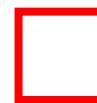
EXAMPLE: BOX FILTER

$f(x, y)$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} w(i, j)$$

$h(x, y)$



EXAMPLE: BOX FILTER

$$f(x, y)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h(x, y)$$

0	10								

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} w(i, j)$$

$$h(x, y) = \sum_{i,j \in \{-1, 0, 1\}} w(i, j) f(x + i, y + j)$$

EXAMPLE: BOX FILTER

$$f(x, y)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h(x, y)$$

0	10	20							

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} w(i, j)$$

$$h(x, y) = \sum_{i,j \in \{-1,0,1\}} w(i, j) f(x + i, y + j)$$

EXAMPLE: BOX FILTER

$$f(x, y)$$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

$$h(x, y)$$

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} w(i, j)$$

$$h(x, y) = \sum_{i,j \in \{-1,0,1\}} w(i, j) f(x + i, y + j)$$

EXAMPLE: BOX FILTER

$$f(x, y)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h(x, y)$$

	0	10	20	30	30				

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$w(i, j)$$

$$h(x, y) = \sum_{i,j \in \{-1, 0, 1\}} w(i, j) f(x + i, y + j)$$

EXAMPLE: BOX FILTER

$$f(x, y)$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h(x, y)$$

	0	10	20	30	30	30	20	10
	0	20	40	60	60	60	40	20
	0	30	60	90	90	90	60	30
	0	30	50	80	80	90	60	30
	0	30	50	80	80	90	60	30
	0	20	30	50	50	60	40	20
	10	20	30	30	30	30	20	10
	10	10	10	0	0	0	0	0

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} w(i, j)$$

$$h(x, y) = \sum_{i,j \in \{-1, 0, 1\}} w(i, j) f(x + i, y + j)$$

EXAMPLE: BOX FILTER

$$f(x, y)$$

SIZE REDUCTION!

	1	1	1
1	1	1	1
9	1	1	1

$$h(x, y)$$

EXAMPLE: BOX FILTER

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect
(remove sharp features)

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} w(i, j)$$



PRACTICE WITH FILTERS



Original

0	0	0
0	1	0
0	0	0

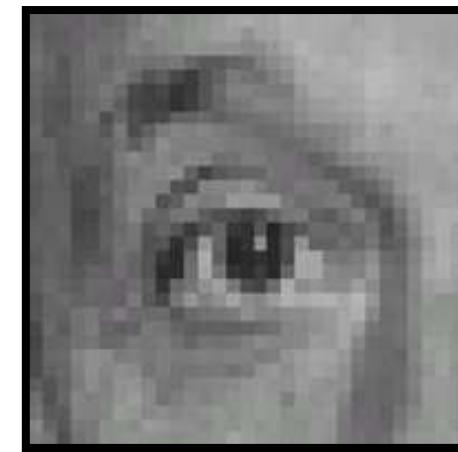
?

PRACTICE WITH FILTERS



Original

0	0	0
0	1	0
0	0	0



Remains the same

PRACTICE WITH FILTERS



Original

0	0	0
0	0	1
0	0	0

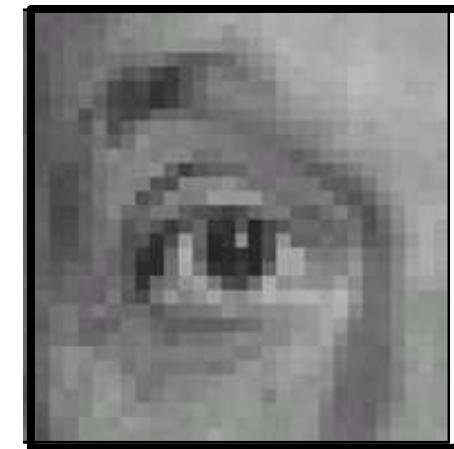
?

PRACTICE WITH FILTERS



Original

0	0	0
0	0	1
0	0	0



Shift 1 pixel to the left

PRACTICE WITH FILTERS



$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

-

$1/9$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

?

PRACTICE WITH FILTERS

Box filter (blurring)



$$\begin{matrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{matrix}$$

- 1/9

$$\begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

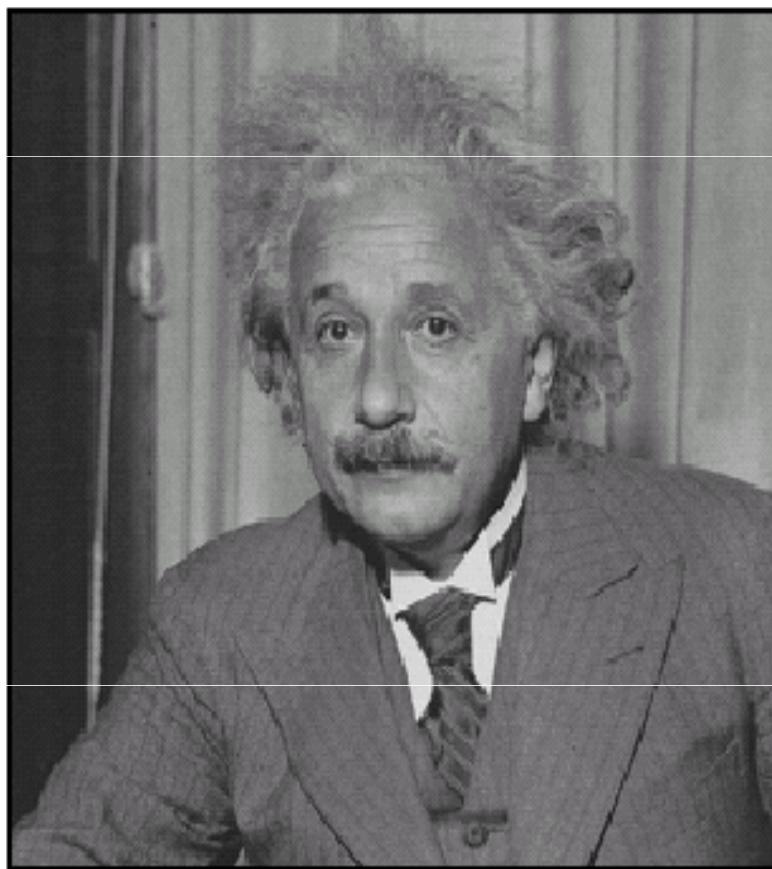


Original

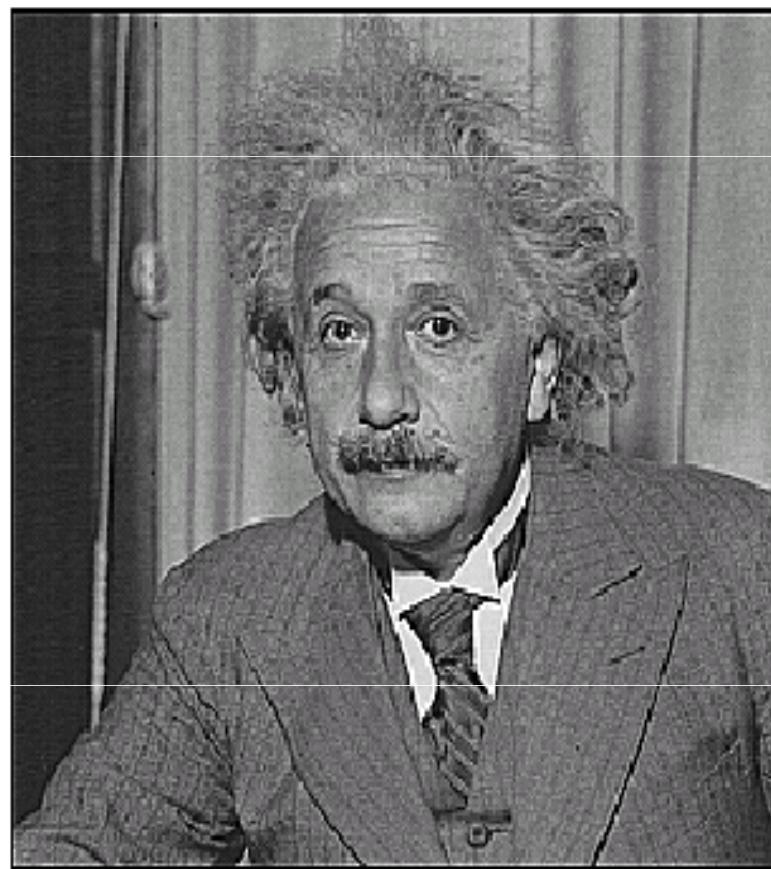
Result = Sharpening filter

- Accentuates differences with local average

SHARPENING

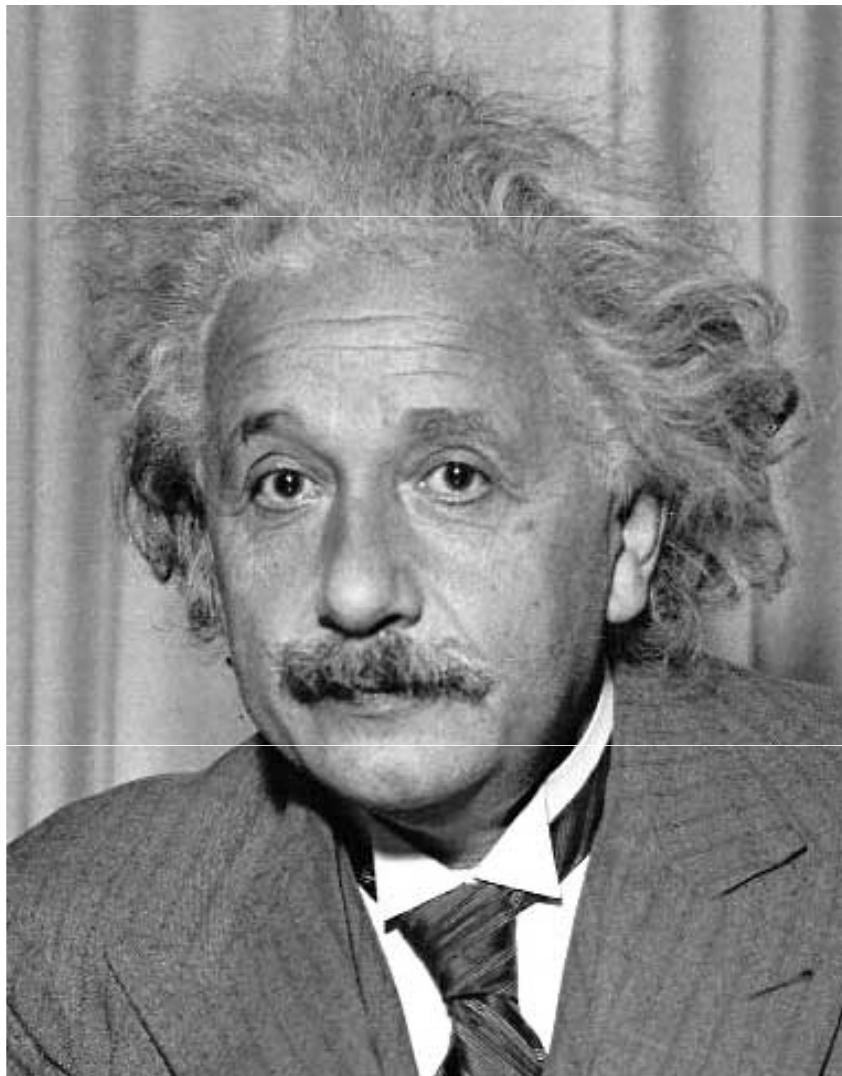


before



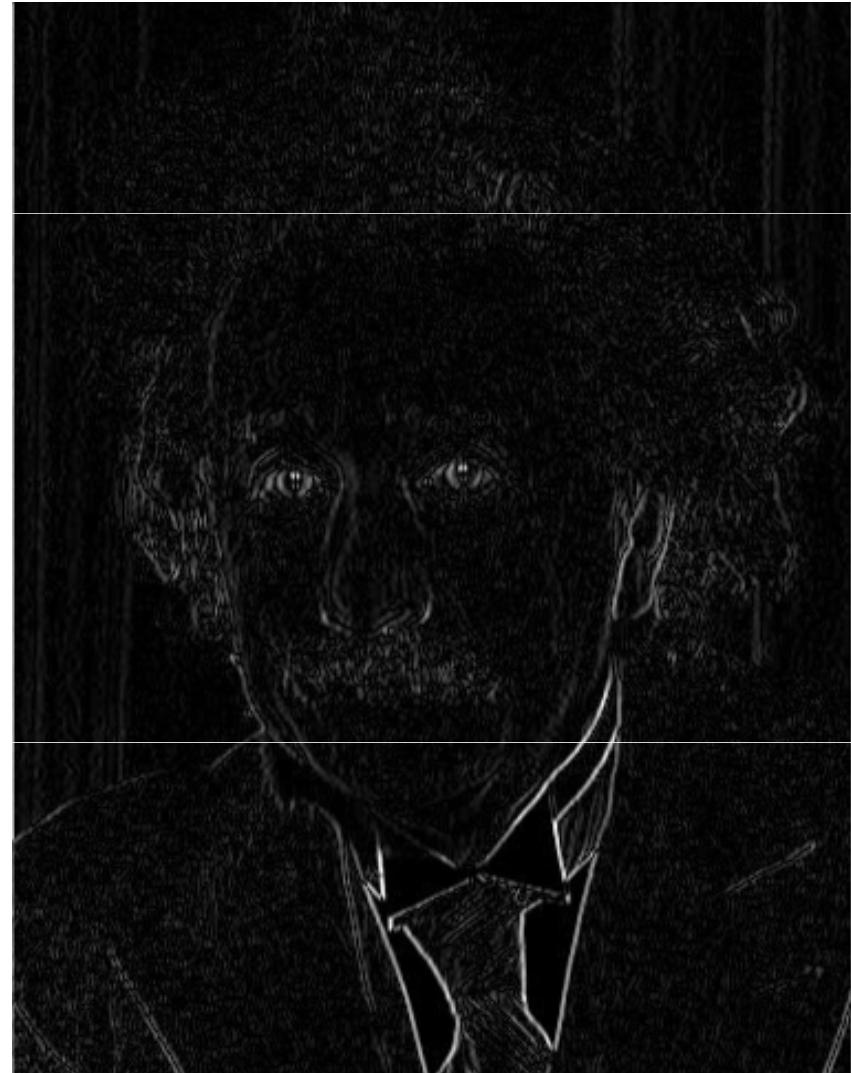
after

PRACTICE WITH FILTERS



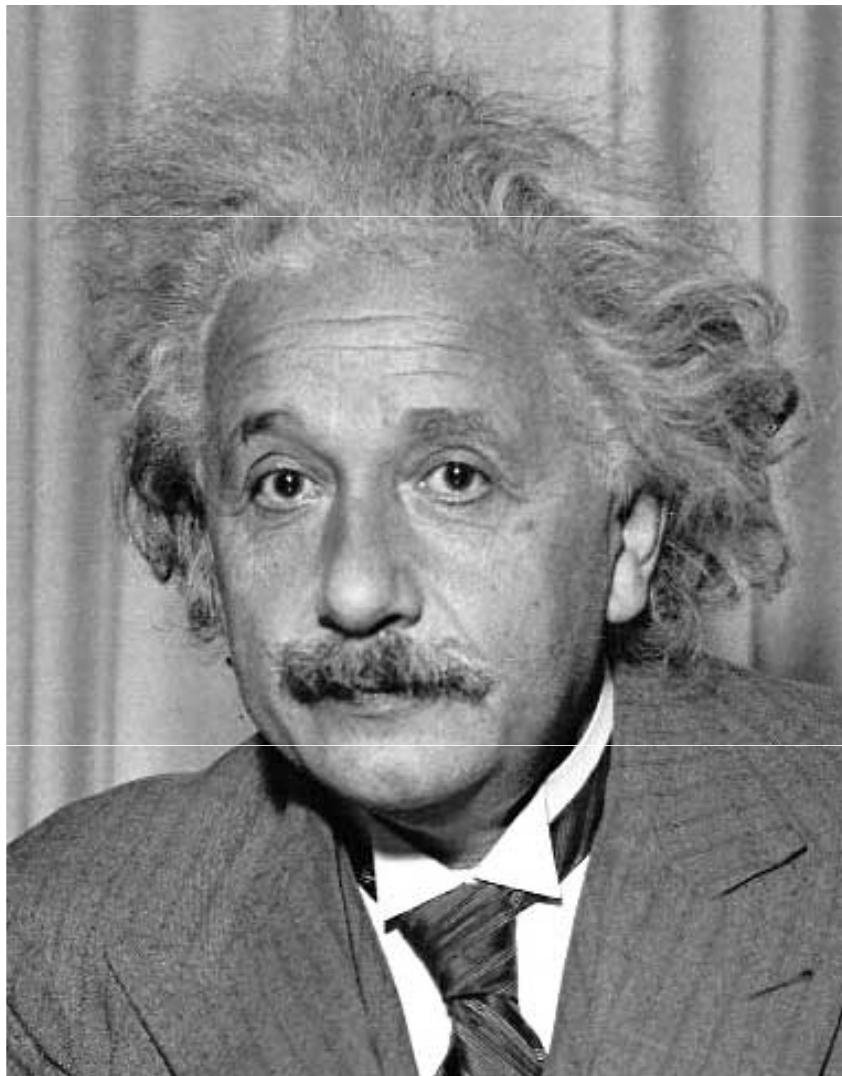
1	0	-1
2	0	-2
1	0	-1

Sobel



Vertical Edge
Horizontal gradient^(absolute value)

PRACTICE WITH FILTERS



1	2	1
0	0	0
-1	-2	-1

Sobel

Vertical gradient



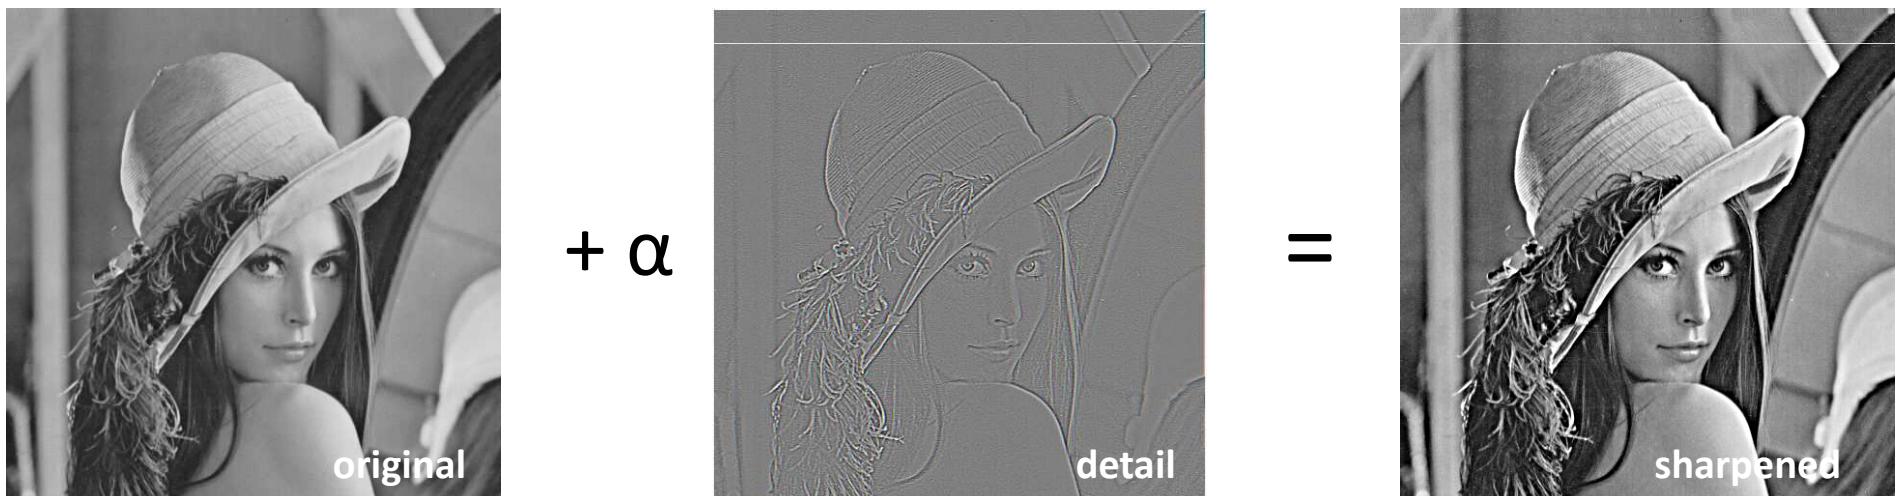
Horizontal Edge
(absolute value)

Sharpening

- What does blurring take away?



Let's add it back:



NOW... FOR THE REAL DEAL

This process of sliding window in general is called **correlation**.

If we multiply the same function it is called **auto-correlation**.

If the functions are different it is also known as **cross-correlation**.

$$w(s, t) \odot f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

The importance of this operation is that it gives a maximum when both functions are the most similar. Thus, the more similar the functions are, the higher the value of the correlation is.

The process of reversing one of the signals and cross-correlate them is the well known **CONVOLUTION**. This is defined as follows,

$$w(s, t) \otimes f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

KEY PROPERTIES OF LINEAR FILTERS

Linearity:

$$\text{filter}(f_1 + f_2) = \text{filter}(f_1) + \text{filter}(f_2)$$

$$\text{filter}(k \cdot f_1 + m \cdot f_2) = k \cdot \text{filter}(f_1) + m \cdot \text{filter}(f_2)$$

Shift invariance: same behavior regardless of pixel location (Use same set of weights at each point)

$$\text{filter}(\text{shift}(f)) = \text{shift}(\text{filter}(f))$$

Is box filtering
linear?



$$w(i, j) = \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

$w(i, j)$

A 3x3 matrix of ones, representing a 3x3 box filter kernel. Below the matrix is the label $w(i, j)$.

Is thresholding linear?



$$g(m, n) = \begin{cases} 255, & f(m, n) > A \\ 0 & otherwise \end{cases}$$

Is box filtering linear?



1	1	1
1	1	1
1	1	1

$$w(i, j)$$

Is thresholding linear?

Thresholding is not linear

$F_1 = 1 > T$ (where $T = 2$)

$F_2 = 2 > T$ (where $T = 2$). Output is 0

$F_1 + F_2 > T$. Output is 1



$$g(m, n) = \begin{cases} 255, & f(m, n) > A \\ 0 & otherwise \end{cases}$$

OTHER PROPERTIES OF LINEAR FILTERS

- Commutative: $a * b = b * a$
 - Conceptually no difference between filter and signal
 -
- Associative: $a * (b * c) = (a * b) * c$
 - Often apply several filters one after another: $((a * b_1) * b_2) * b_3$
 - This is equivalent to applying one filter: $a * (b_1 * b_2 * b_3)$
- Distributes over addition: $a * (b + c) = (a * b) + (a * c)$
- Scalars factor out: $ka * b = a * kb = k(a * b)$
- Identity: unit impulse $e = [0, 0, 1, 0, 0]$,
- $a * e = a$

Correlation

This process of sliding window in general is called **correlation**.

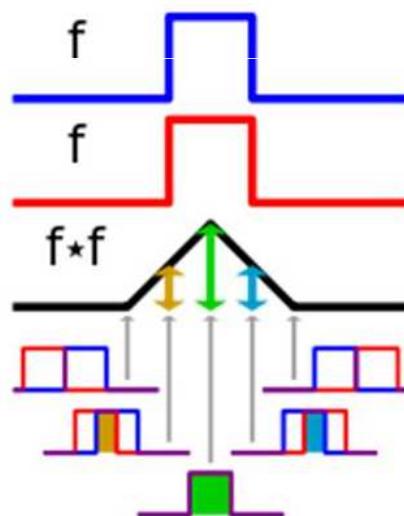
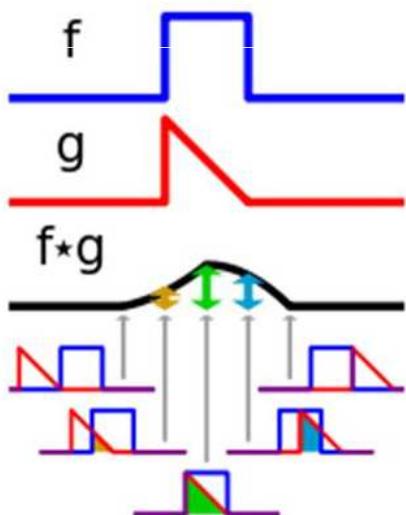
If we multiply the same function it is called **auto-correlation**.

If the functions are different it is also known as **cross-correlation**.

$$w(s, t) \odot f(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

Cross-correlation and autocorrelation

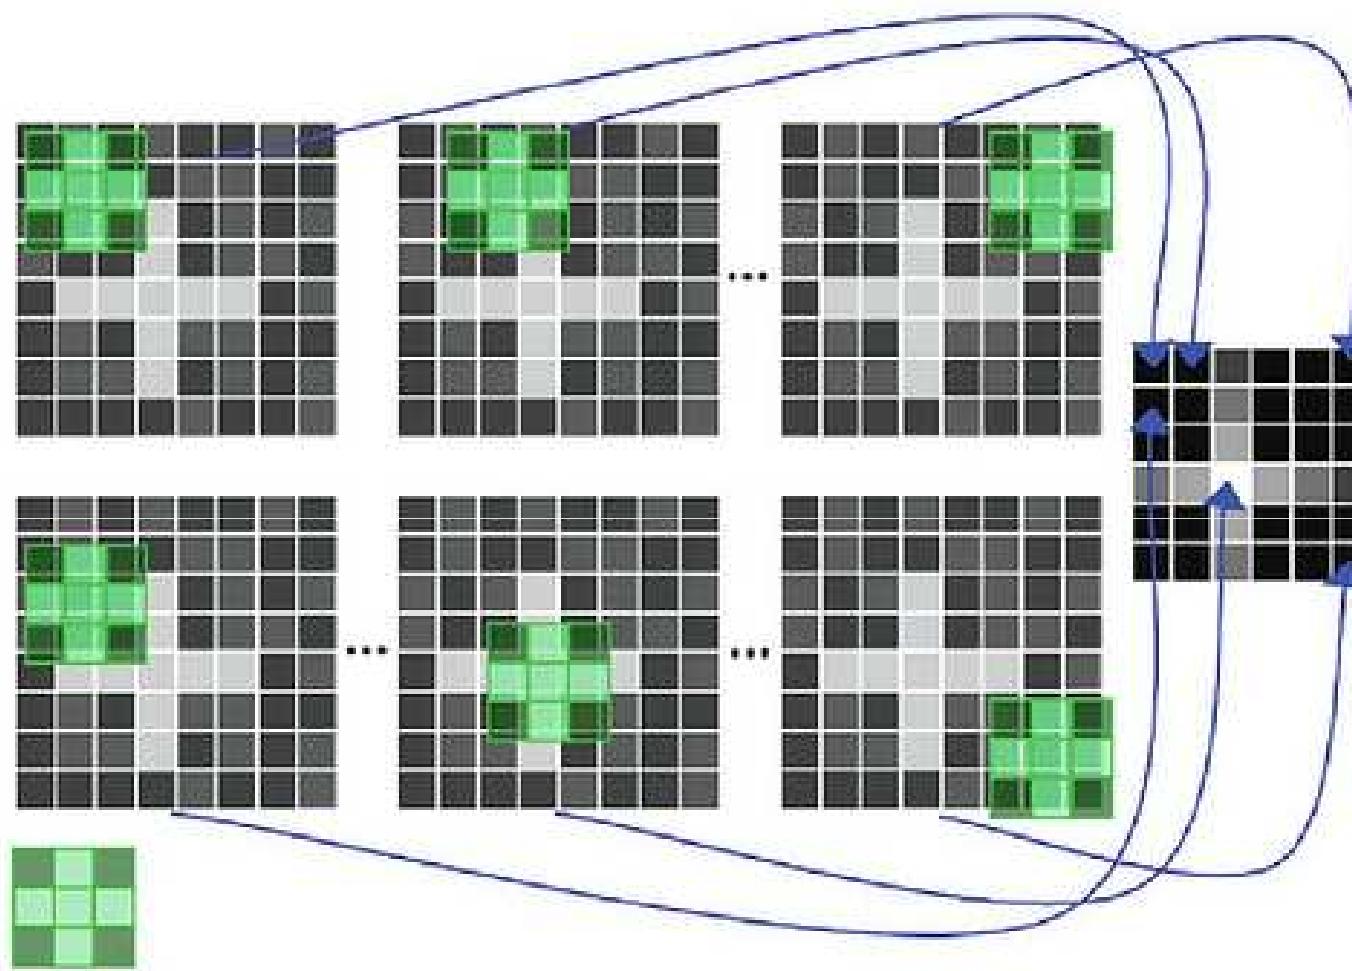
$$(f \times g)(t) = \sum_{\tau} f(\tau)g(t+\tau) \quad (f * f)(t) = \sum_{\tau} f(\tau)f(t-\tau)$$



The importance of this operation is that it gives a maximum when both functions are the most similar. Thus, the more similar the functions are, the higher the value

Correlation

Example in 2D:



Cross-correlation

Let F be the image, H be the kernel (of size $2k+1 \times 2k+1$), and G be the output image

- Can think of as a “dot product” between local neighborhood and kernel for each pixel

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v]F[i + u, j + v]$$

This is called a **cross-correlation** operation:

$$G = H \otimes F$$

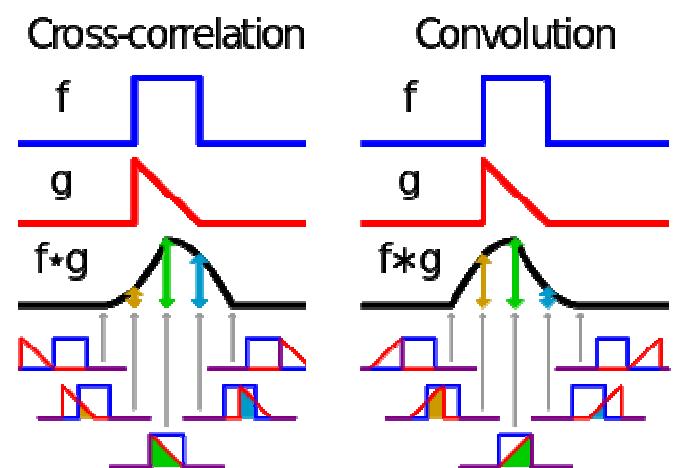
Convolution

- Same as cross-correlation, except that the kernel is “flipped” (horizontally and vertically)

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

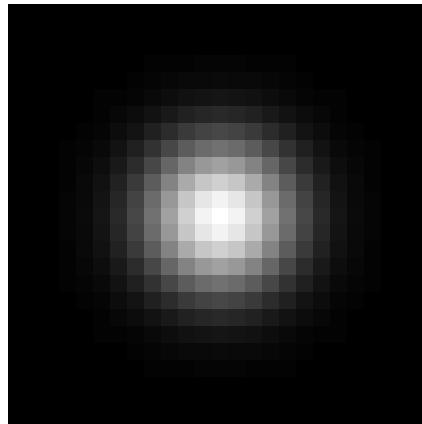
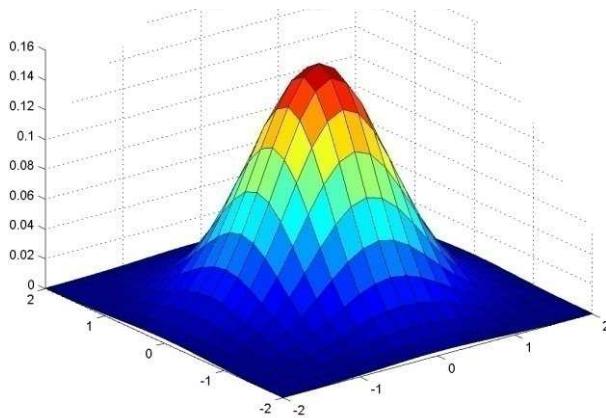
Convolution is **commutative** and **associative**

$$G = H * F$$



AN IMPORTANT FILTER: GAUSSIAN

- Weight contributions of neighboring pixels by nearness

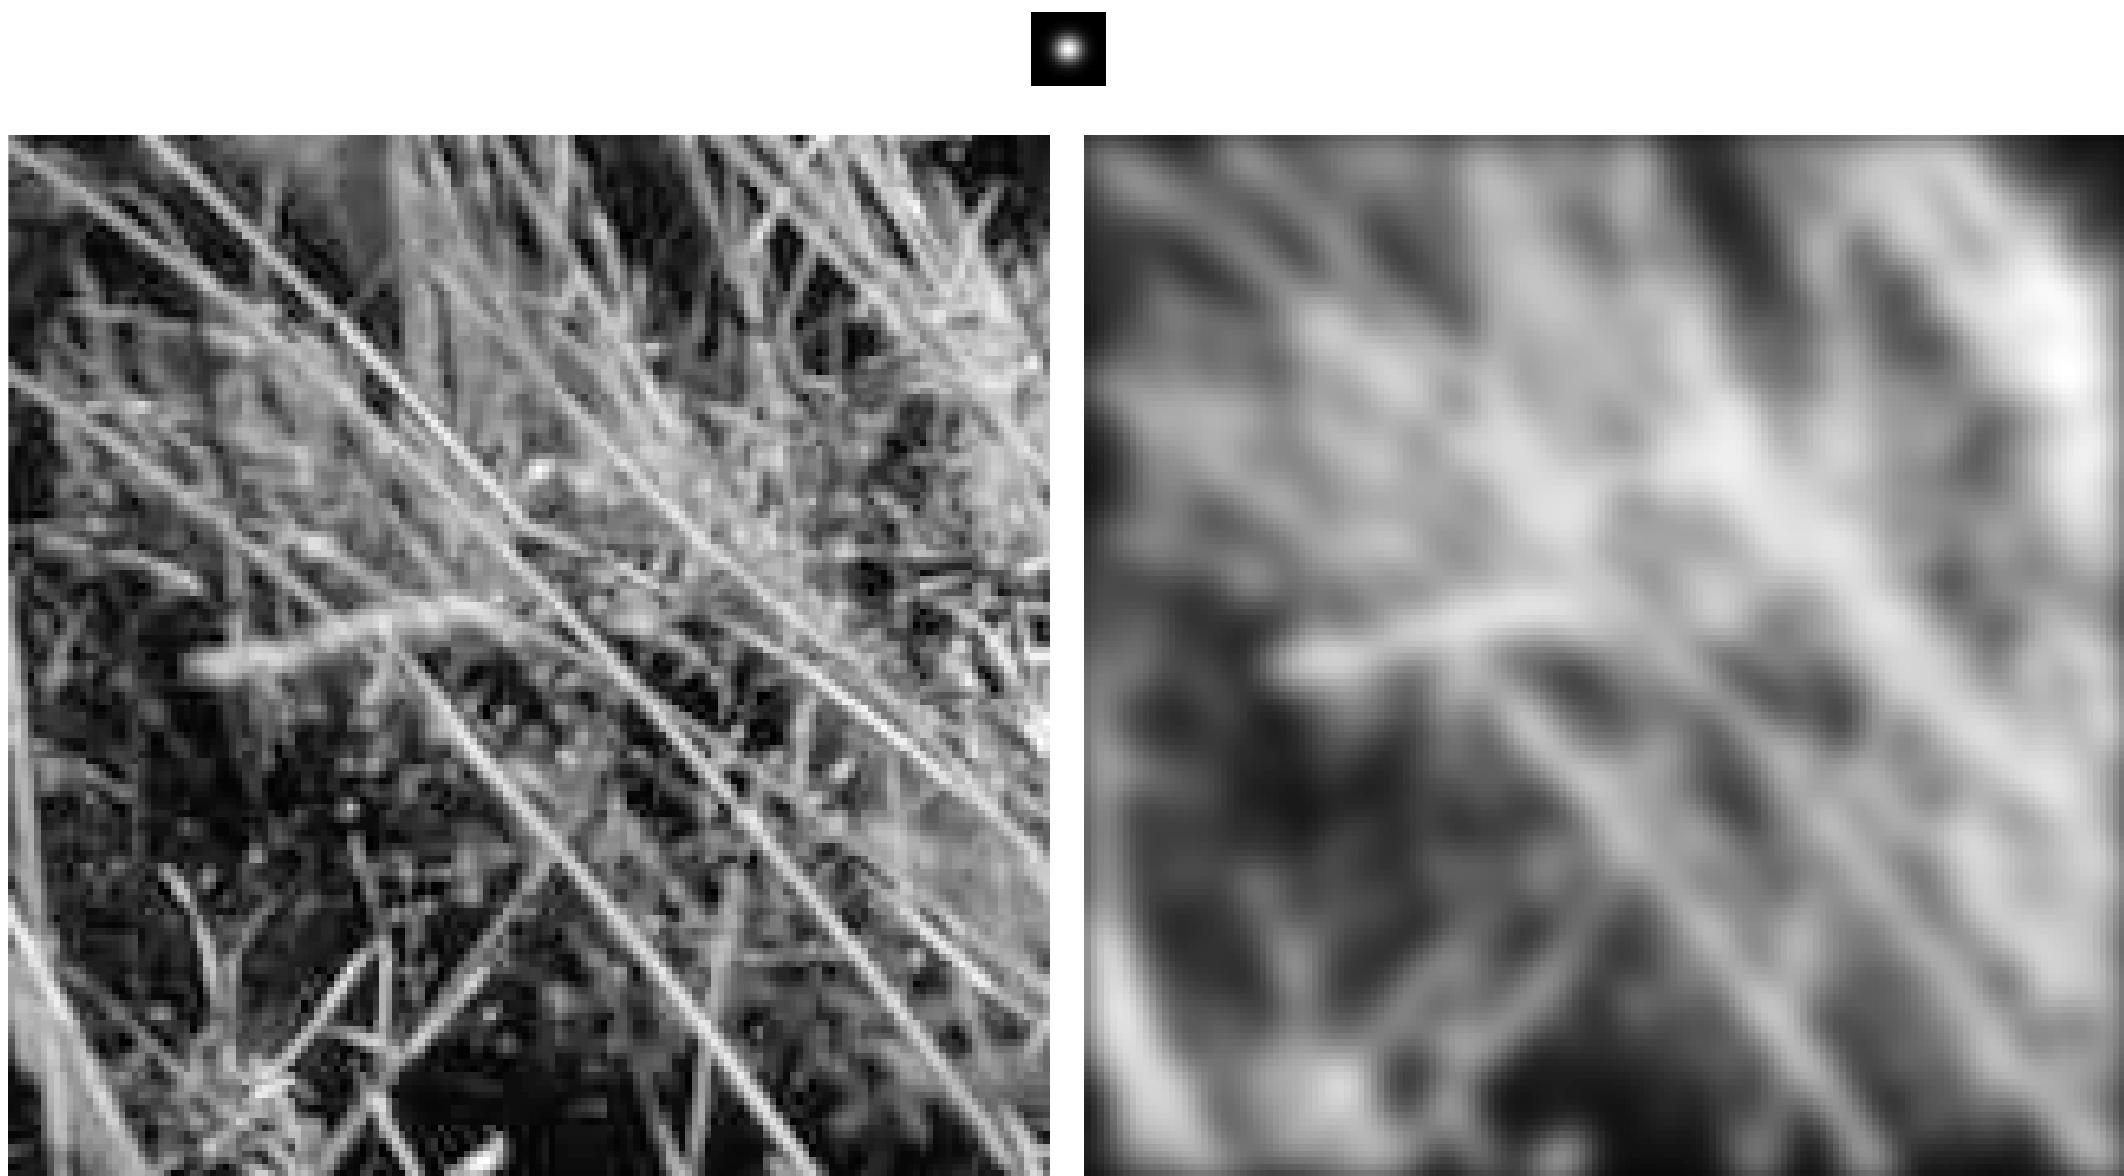


0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

$5 \times 5, \sigma = 1$

$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

SMOOTHING WITH A GAUSSIAN FILTER



SMOOTHING WITH A BOX FILTER



GAUSSIAN FILTER PROPERTIES

- Remove “high-frequency” components from the image (low-pass filter)
 - Images become more smooth
- Convolution with self is another Gaussian
 - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
 - Convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- *Separable* kernel
 - Factors into product of two 1D Gaussians

GAUSSIAN FILTER PROPERTIES

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

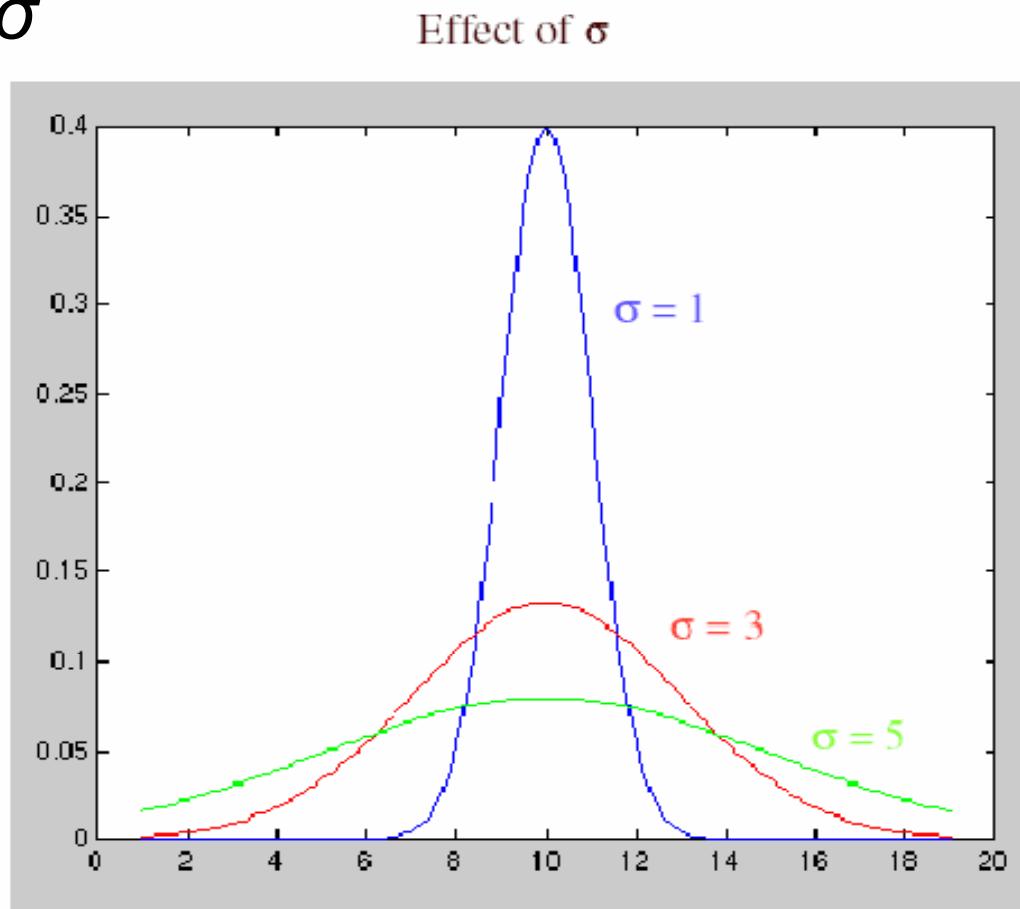
The 2D Gaussian can be expressed as the product of two functions, one a function of x and the other a function of y

In this case, the two functions are the (identical) 1D Gaussian

PRACTICAL MATTERS

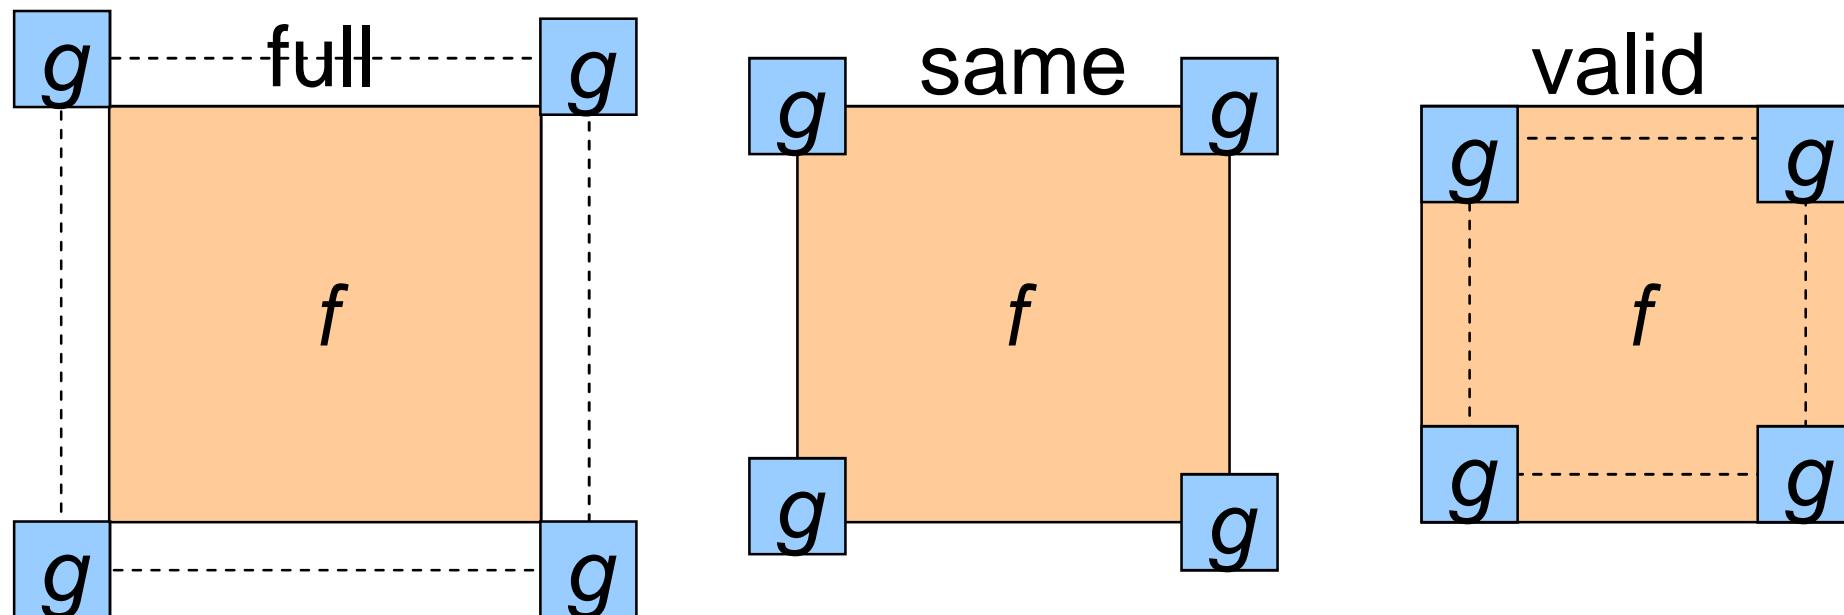
How big should the filter be?

- Values at edges should be near zero
- Rule of thumb for Gaussian: set filter half-width to about 3σ



PRACTICAL MATTERS

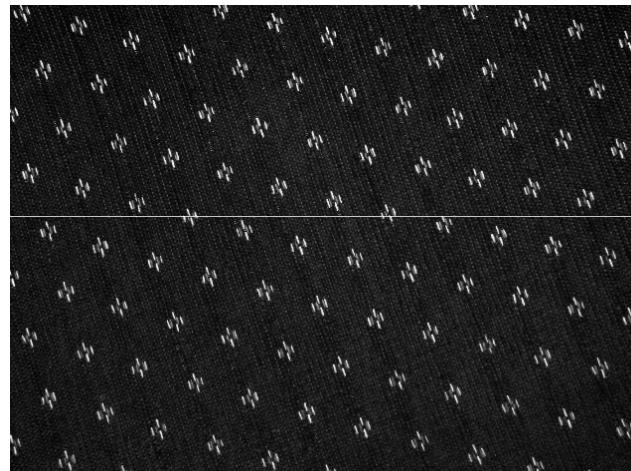
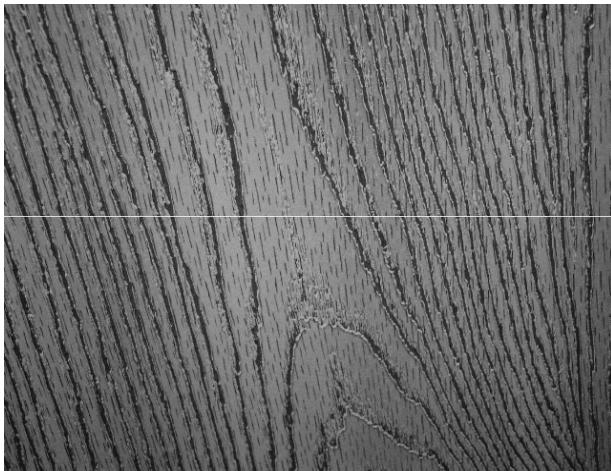
- What is the size of the output?
- MATLAB: `filter2(g, f, shape)`
 - `shape = 'full'`: output size is sum of sizes of f and g
 - `shape = 'same'`: output size is same as f
 - `shape = 'valid'`: output size is difference of sizes of f and g



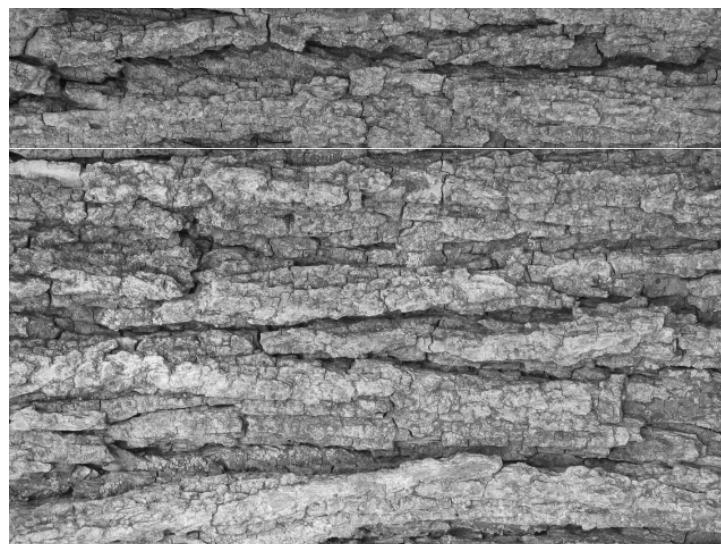
APPLICATION OF FILTERS: TEXTURES



TEXTURES AND MATERIALS



TEXTURES AND ORIENTATIONS



TEXTURES AND SCALES



WHAT IS A TEXTURE?

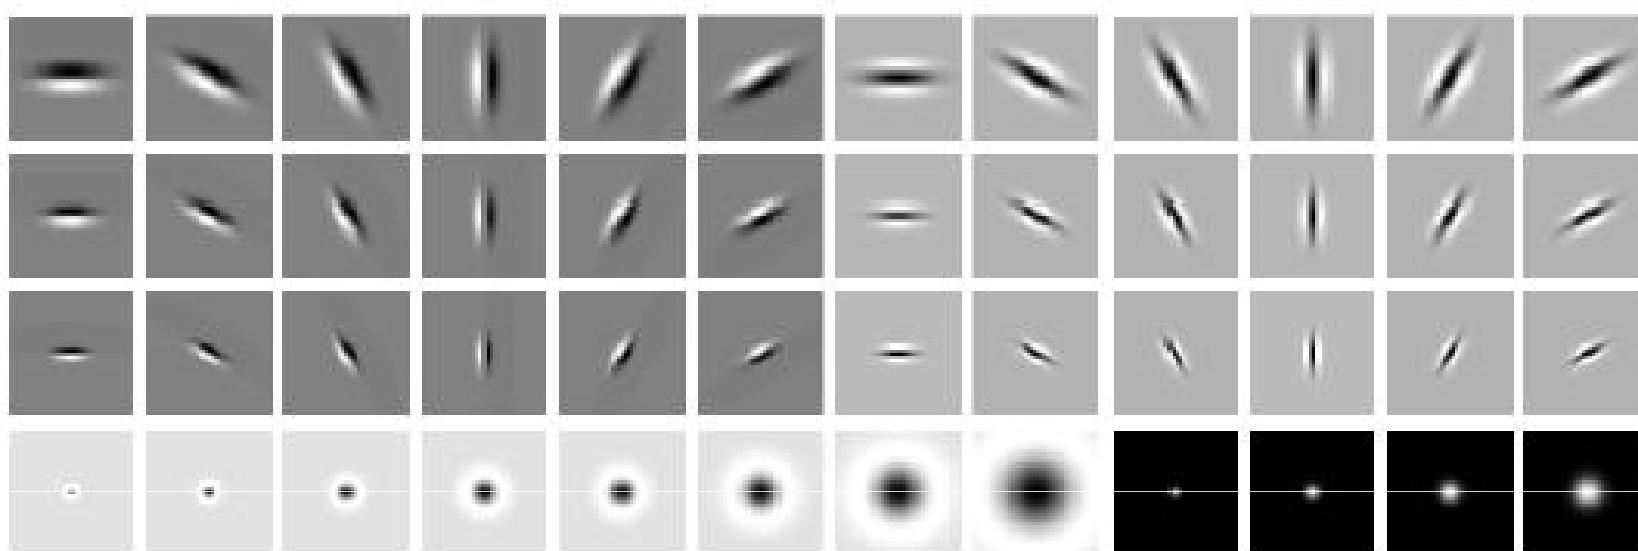
Regular or stochastic patterns caused by bumps, grooves, and/or markings

HOW CAN WE REPRESENT TEXTURE?

Compute responses of blobs and edges at various orientations and scales

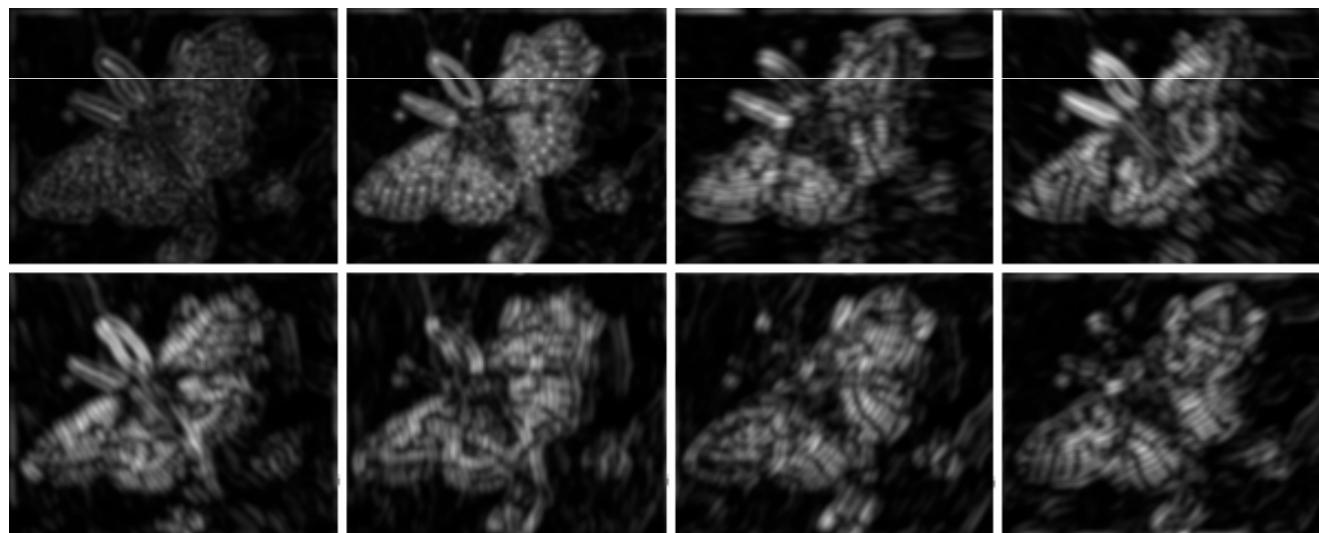
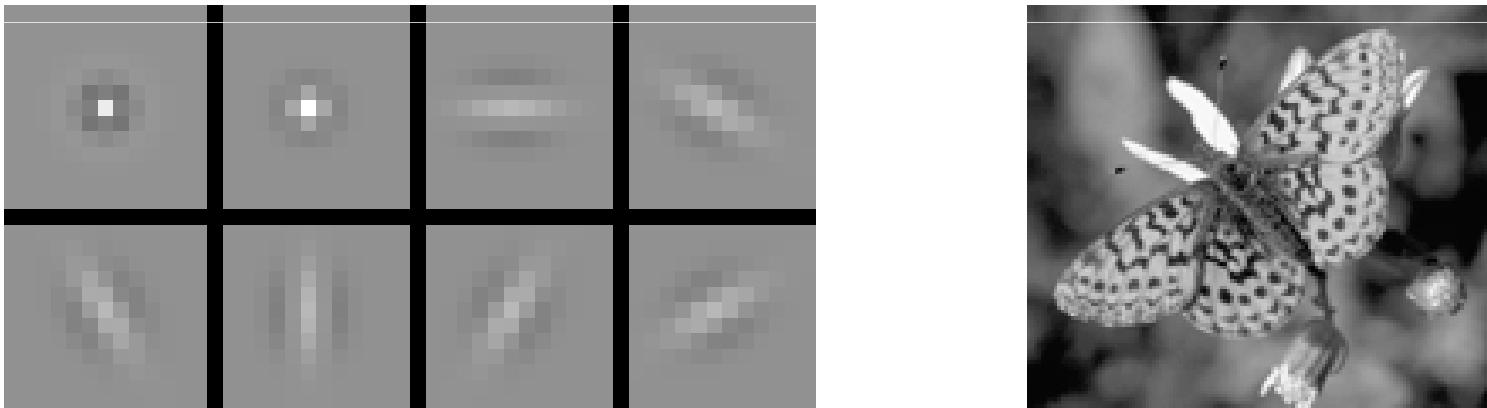
OVERCOMPLETE REPRESENTATION: FILTER BANKS

LM Filter Bank

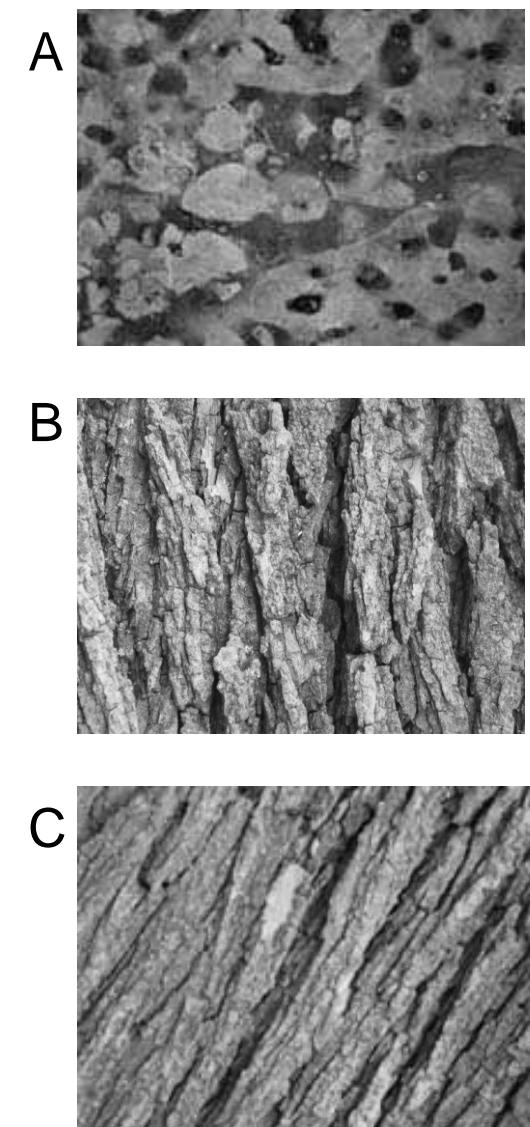
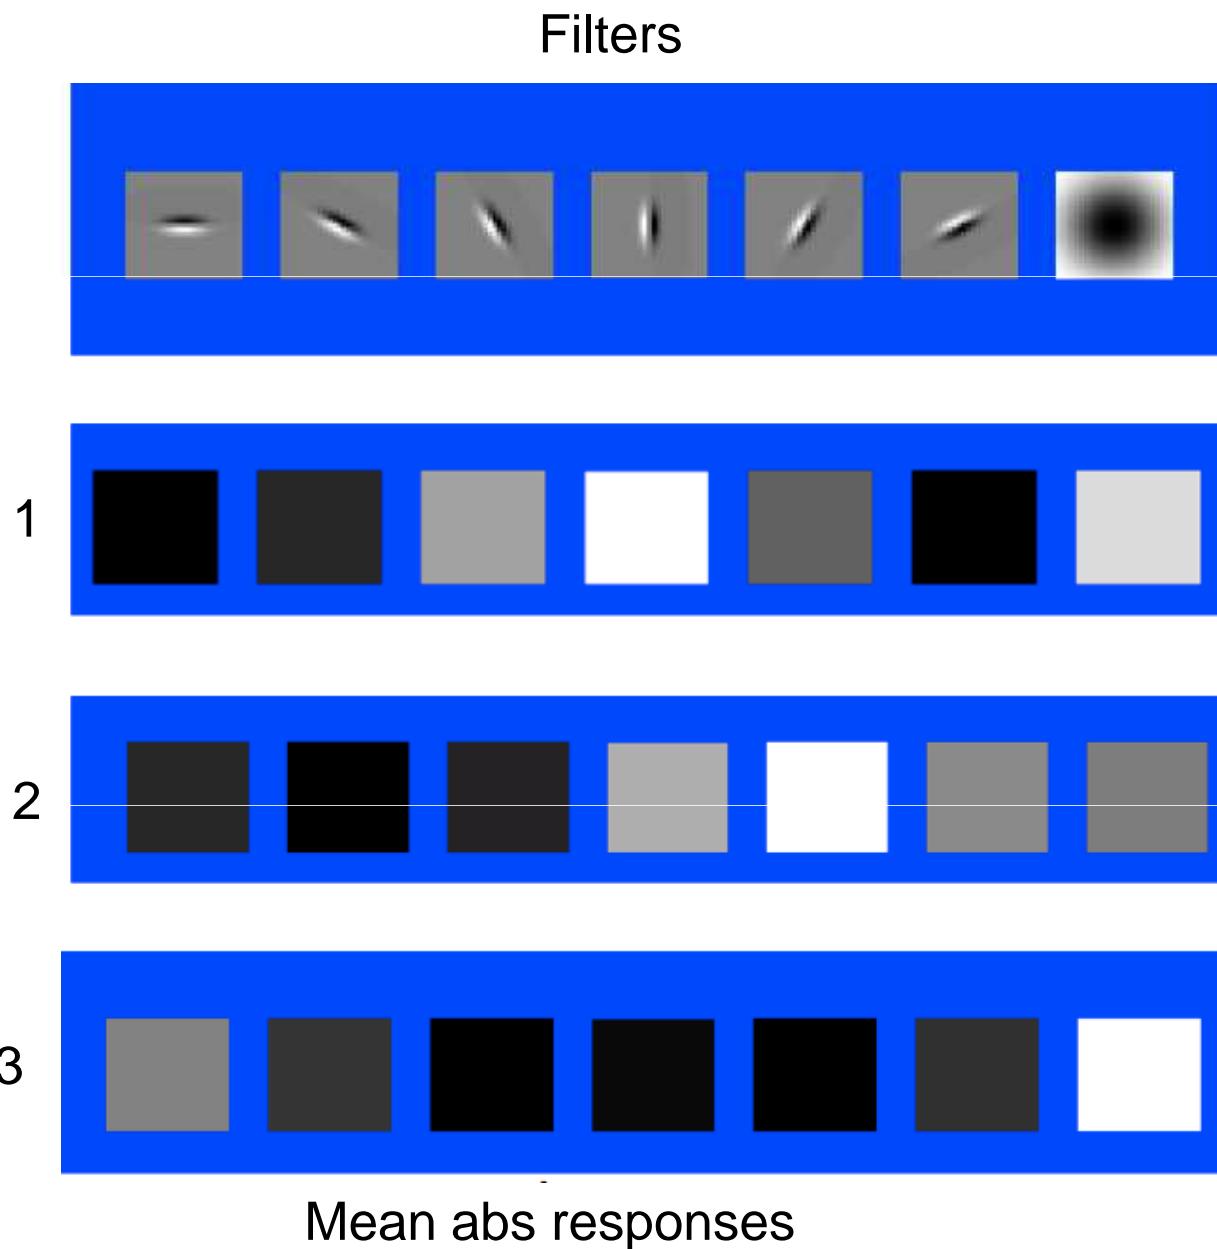


HOW TO USE IT...

- Process image with each filter and keep responses (or squared/abs responses)

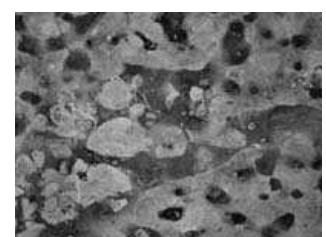
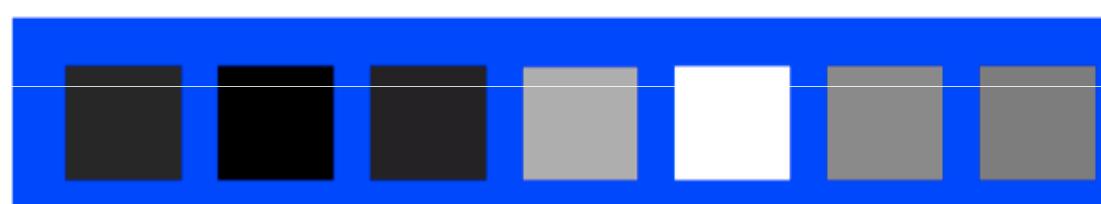
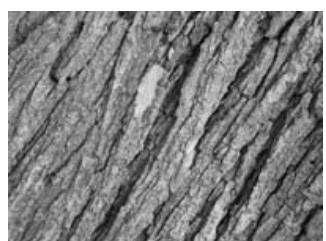
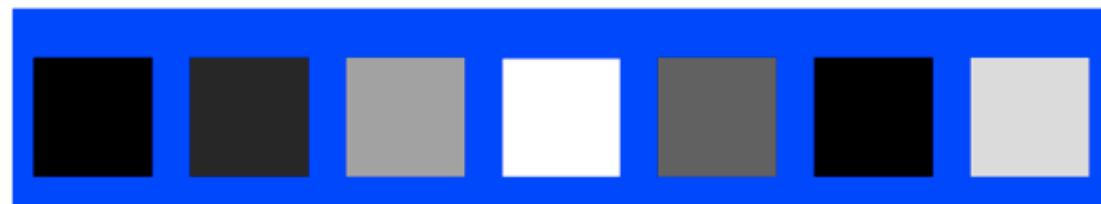
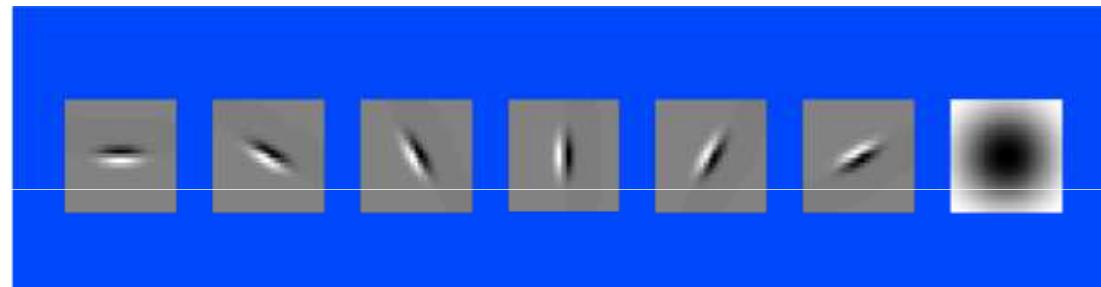


MATCH TEXTURE WITH RESPONSE



MATCH TEXTURE WITH RESPONSE

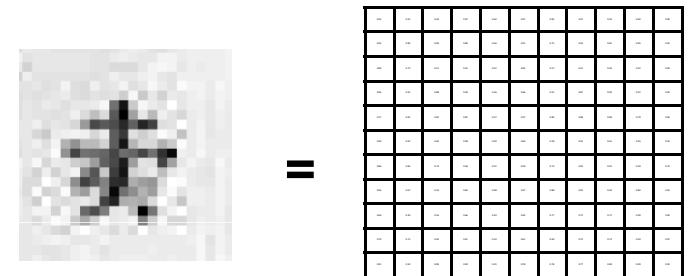
Filters



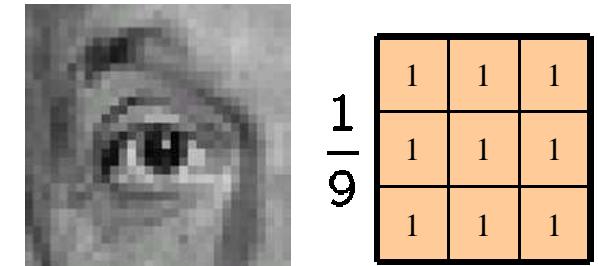
Mean abs responses

TO REMEMBER

- Image is a matrix of numbers


$$\text{Image} = \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline \end{array}$$

- Linear filtering is sum of dot product at each position
 - Can smooth, sharpen, translate (among many other uses)


$$\frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

- Be aware of details for filter size, extrapolation, cropping



REVIEW QUESTIONS

1. Write down a 3x3 filter that returns a positive value if the average value of the 4-adjacent neighbors is less than the center and a negative value otherwise
1. Write down a filter that will compute the gradient in the x-direction:

$$\text{grad}_x(y, x) = \text{im}(y, x+1) - \text{im}(y, x) \text{ for each } x, y$$

REVIEW QUESTIONS

3. Fill in the blanks:

Filtering Operator

a) $\underline{\quad} = D * B$

b) $A = \underline{\quad} * \underline{\quad}$

c) $F = D * \underline{\quad}$

d) $\underline{\quad} = D * D$

